



Containers on AWS

AWS KRUG 컨테이너 소모임

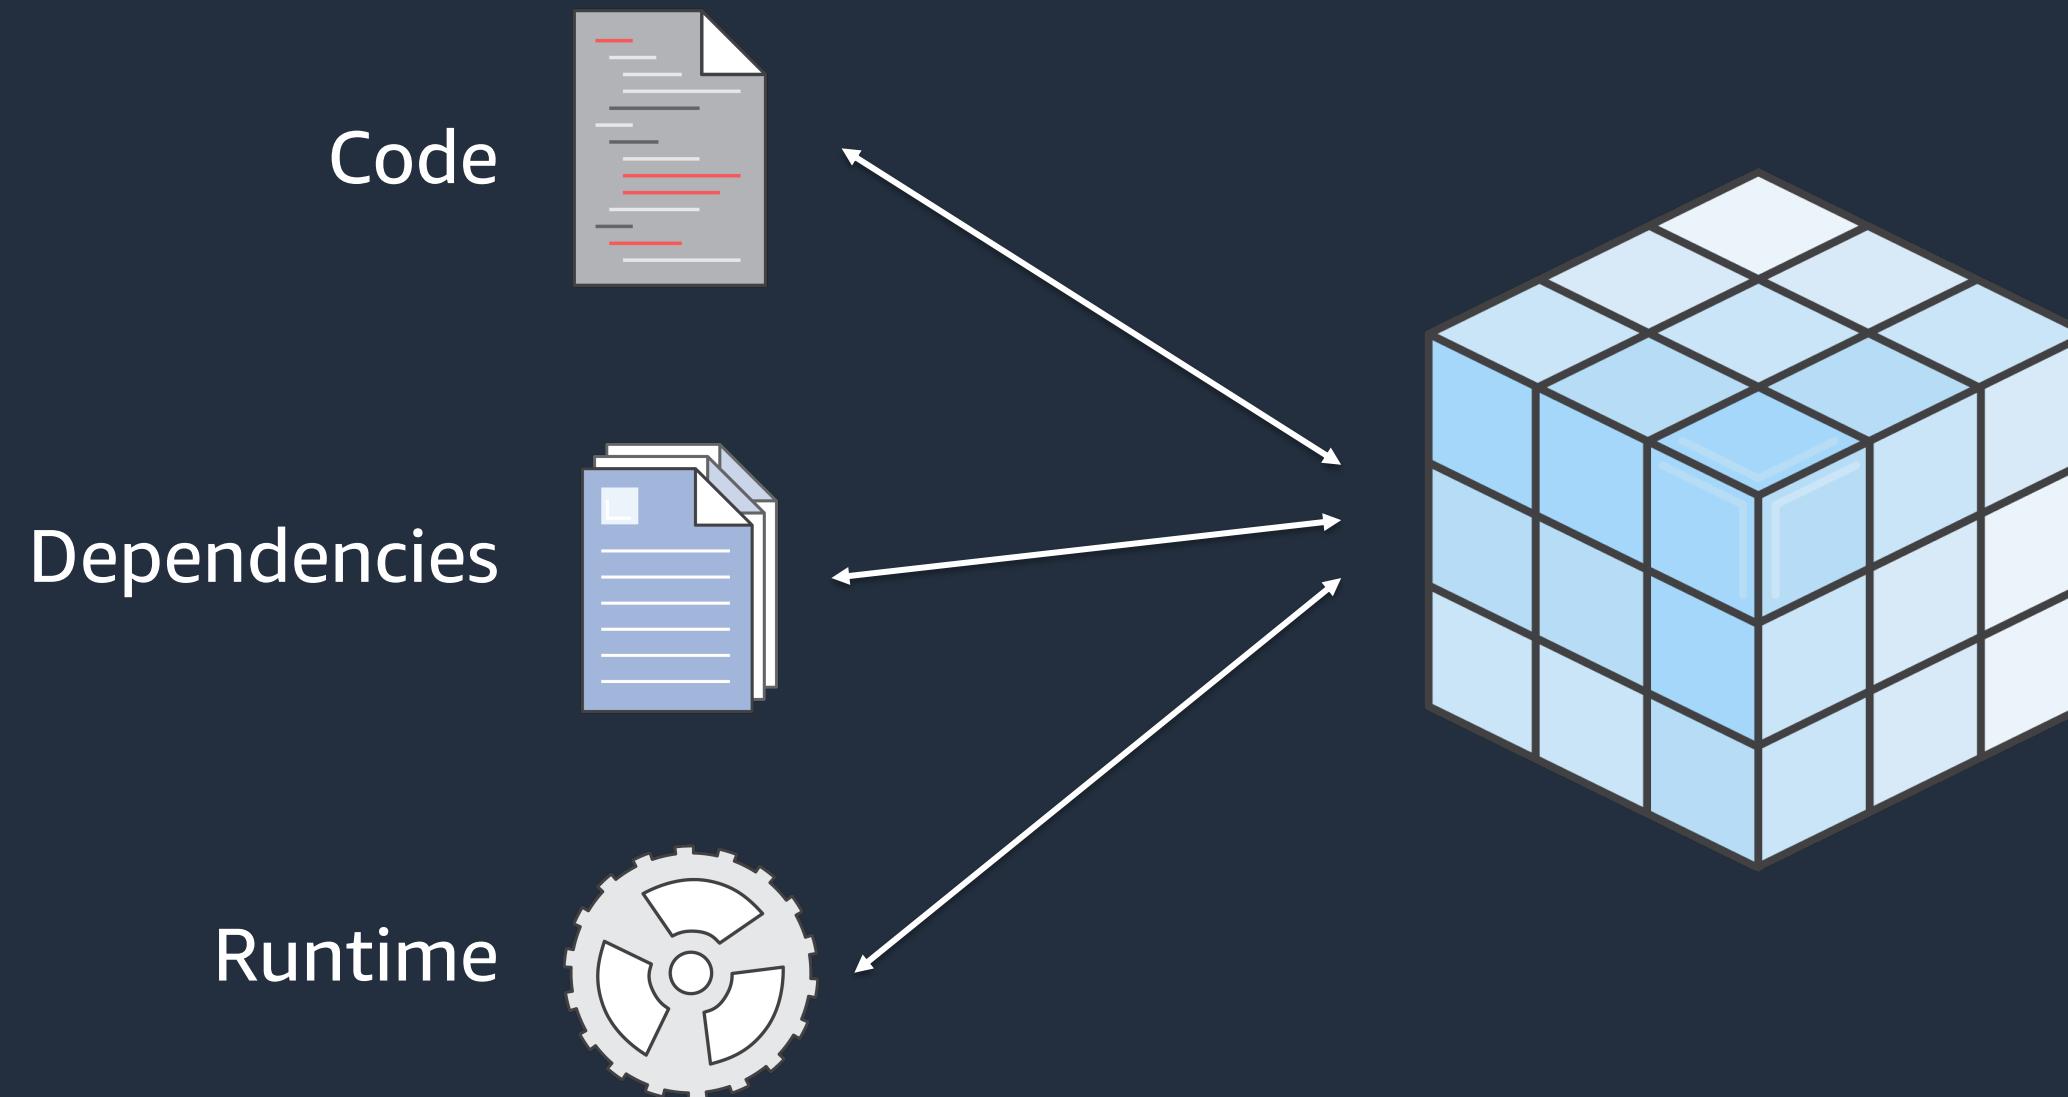
김지민 솔루션즈 아키텍트, AWS



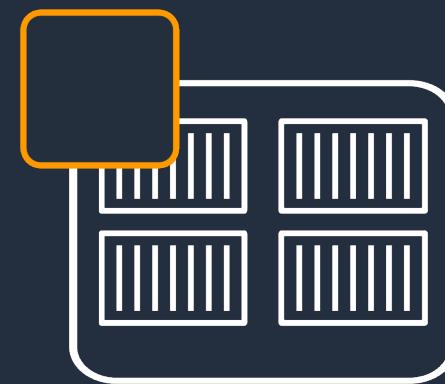
Agenda

1. 컨테이너가 대세인 이유?
2. (너무) 다양한 AWS 컨테이너 서비스들
3. Amazon EKS, ECS, Fargate
4. 컨테이너 애플리케이션 운영하기 - 로깅, 모니터링
5. 쉽게 컨테이너 시작하기

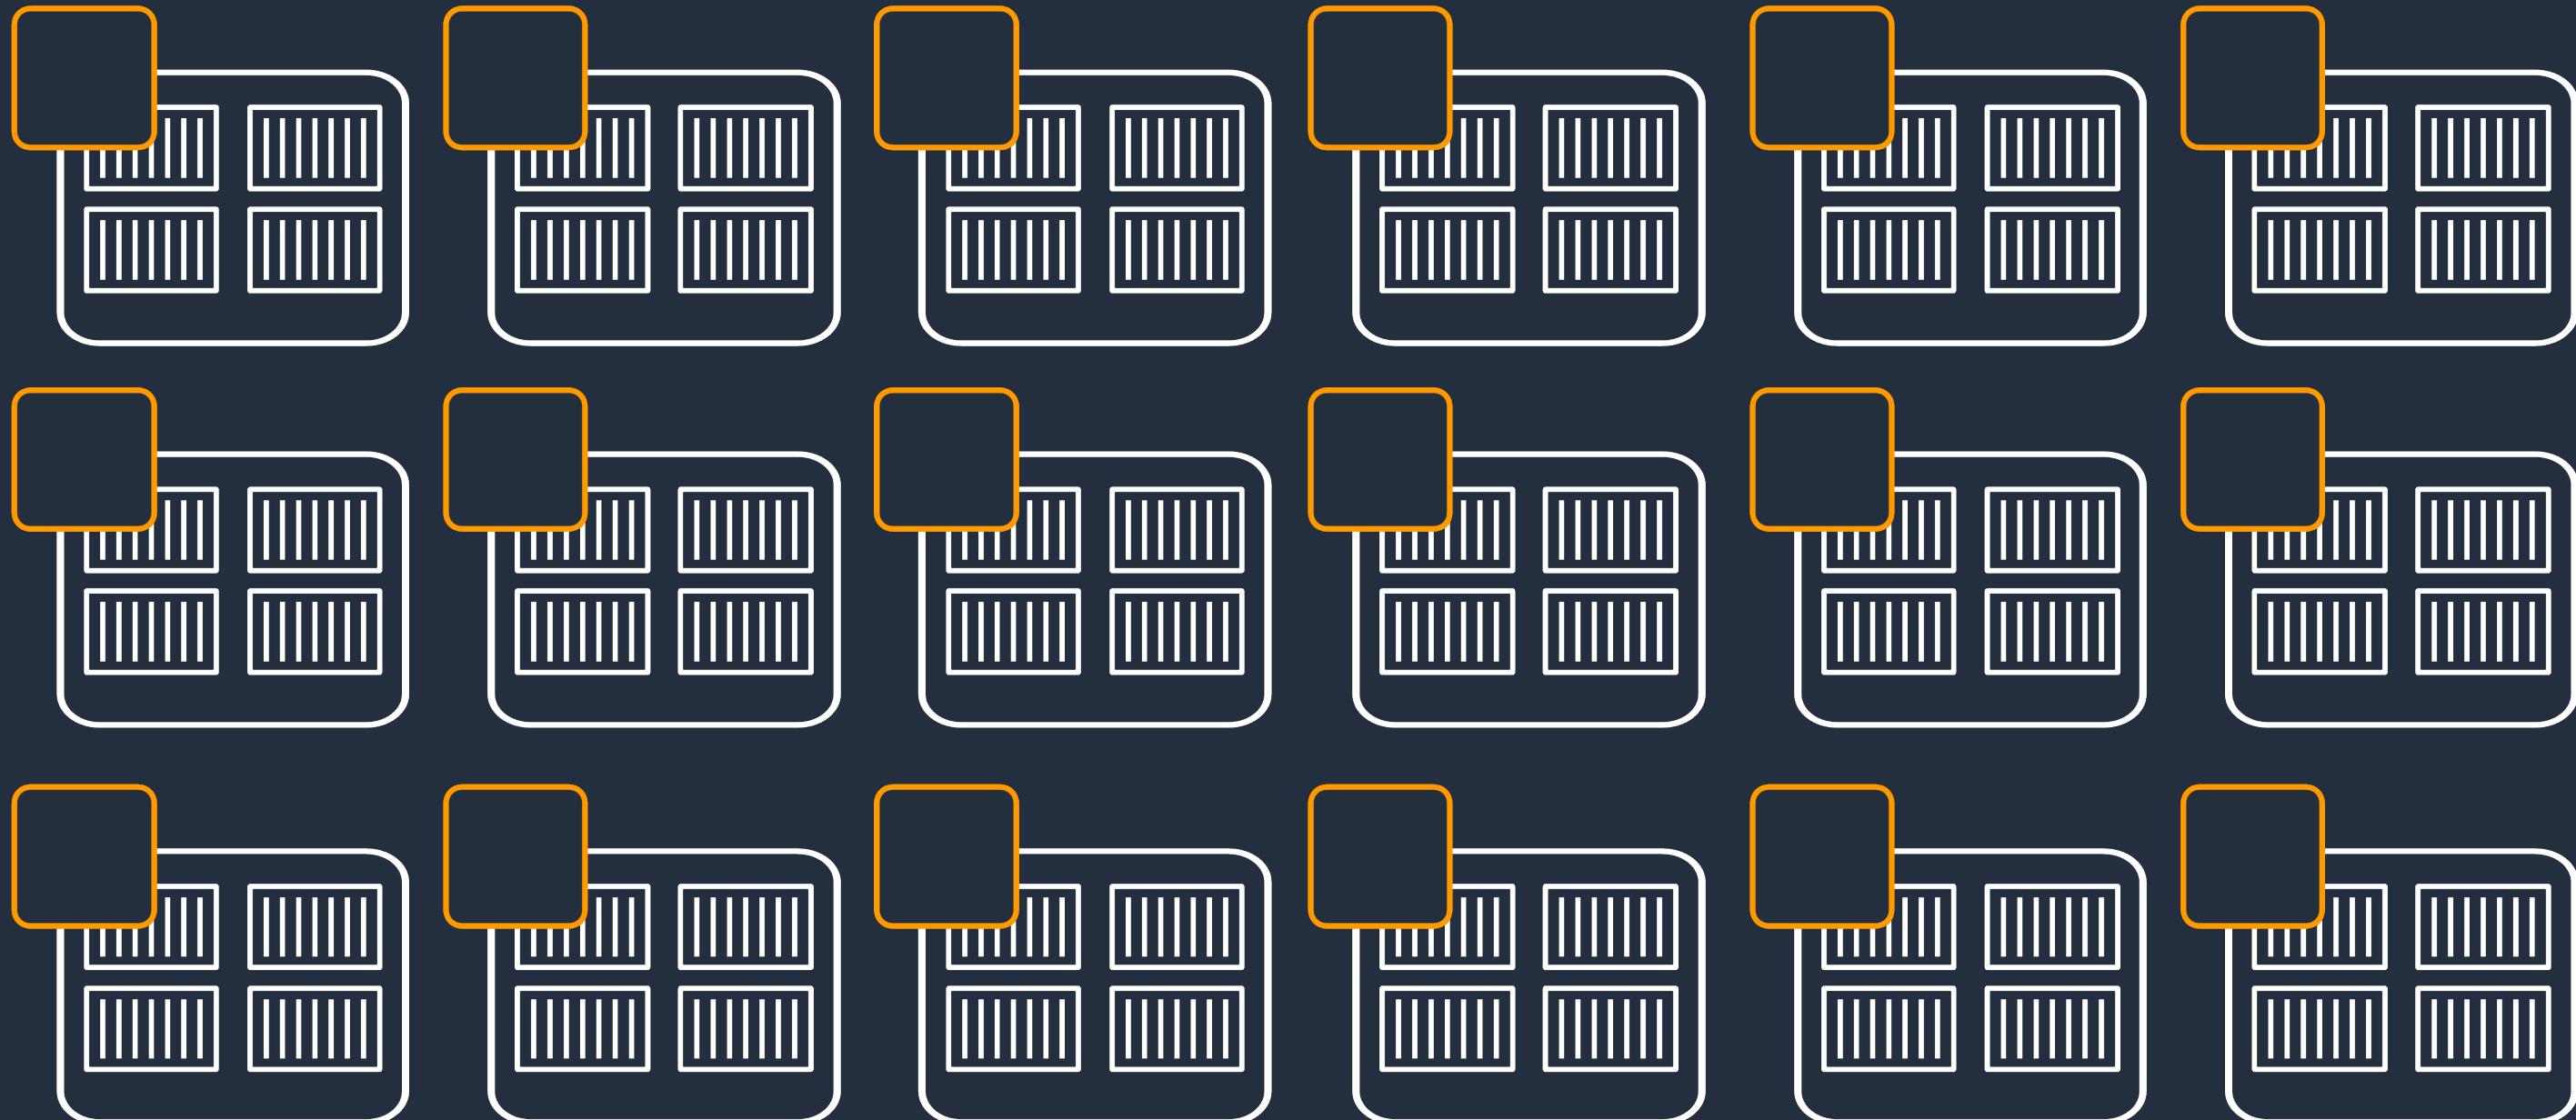
컨테이너란?



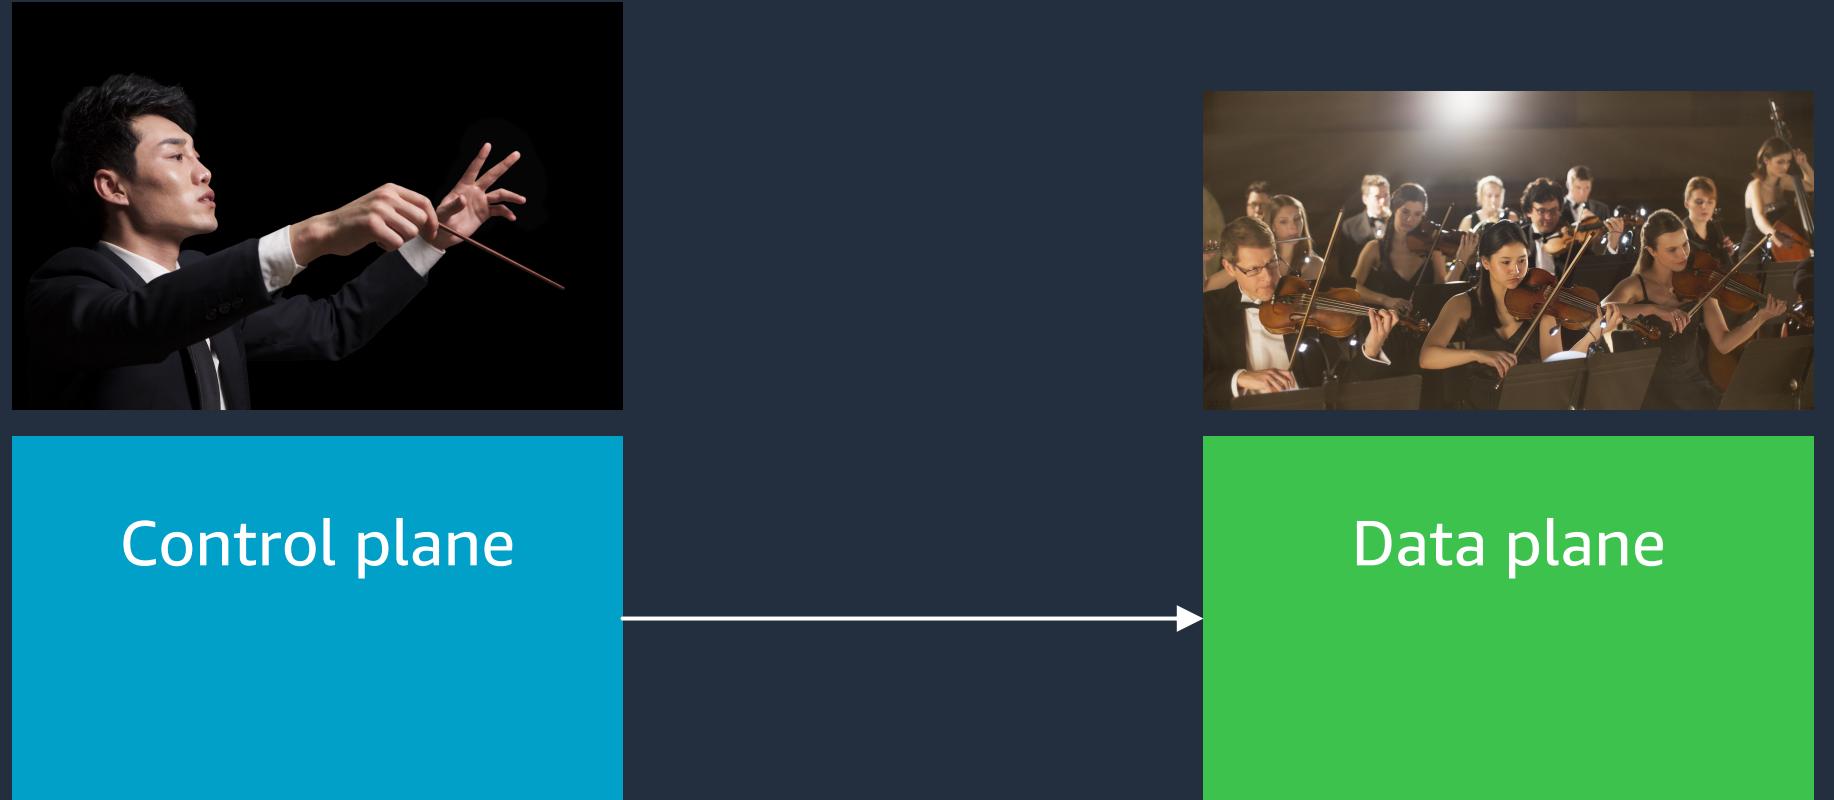
하나의 컨테이너로 시작



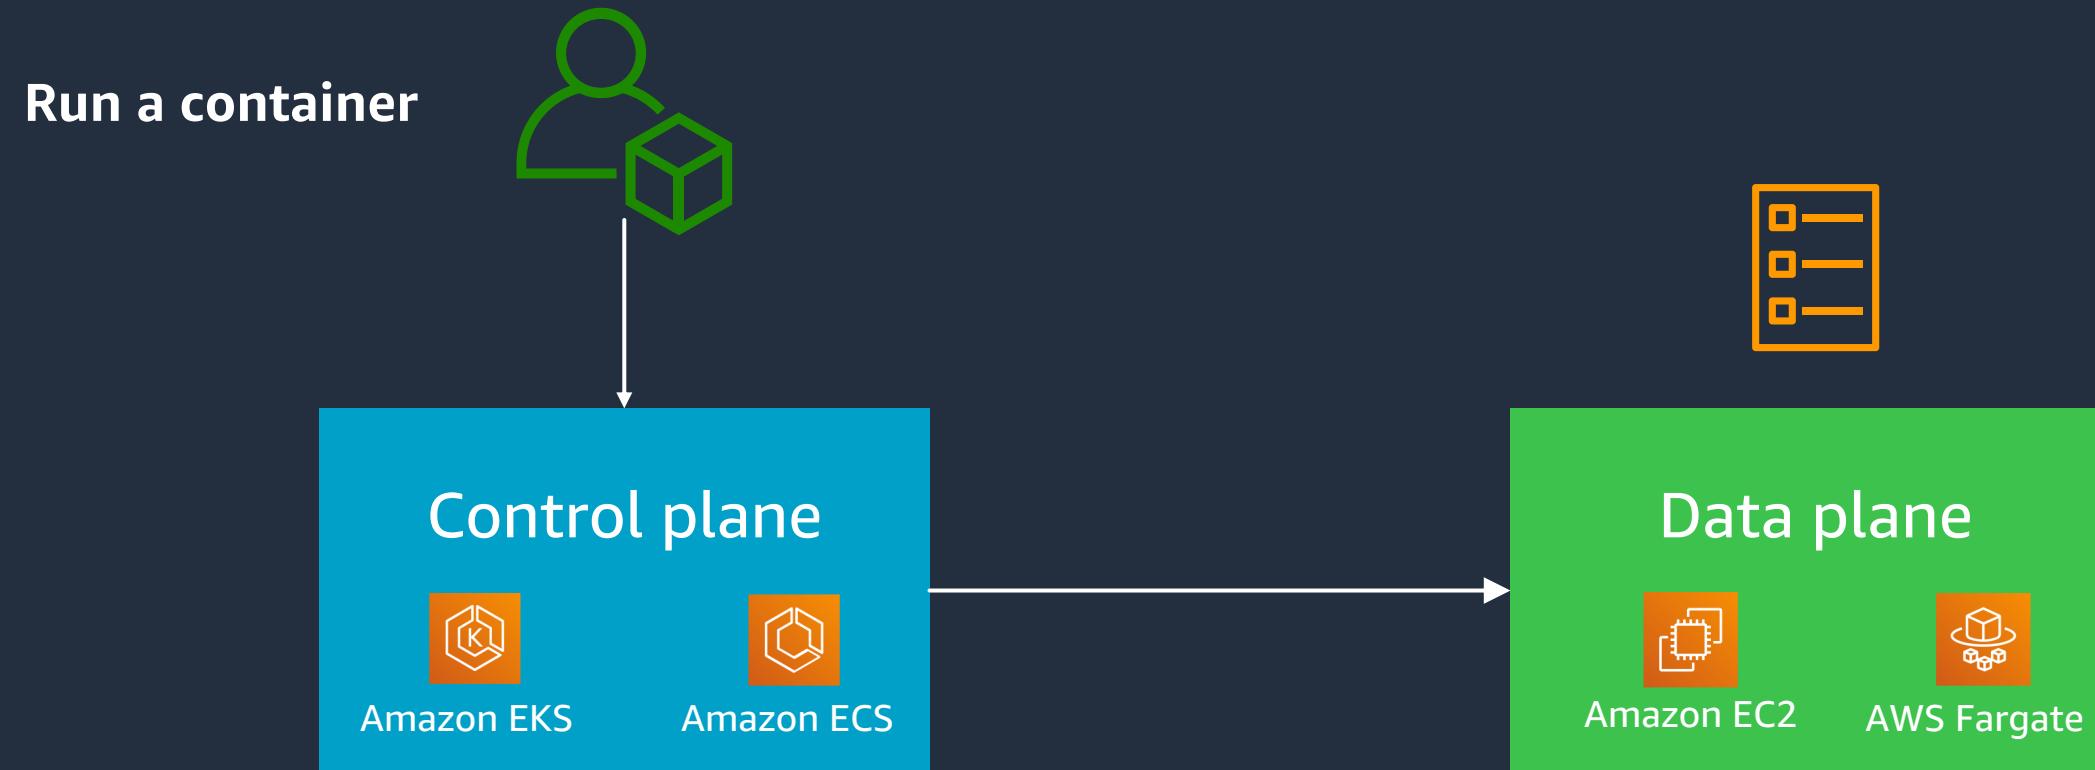
컨테이너가 아주 많아지게 된다면?



컨테이너 오케스트레이터



컨테이너 오케스트레이터



사용자가 원하는 상태(**Desired State**)로 동작하도록 관리(**Schedule**)하는 것

AWS container services landscape

APPLICATION NETWORKING

Service discovery and service mesh



AWS Cloud Map



AWS App Mesh

MANAGEMENT

Deployment, scheduling,
scaling, and management
of containerized applications



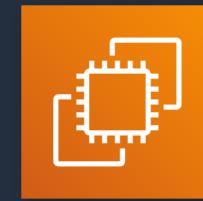
Amazon Elastic
Container Service
(Amazon ECS)



Amazon Elastic
Kubernetes Service
(Amazon EKS)

HOSTING

Where the containers run



Amazon Elastic
Compute Cloud
(Amazon EC2)



AWS Fargate

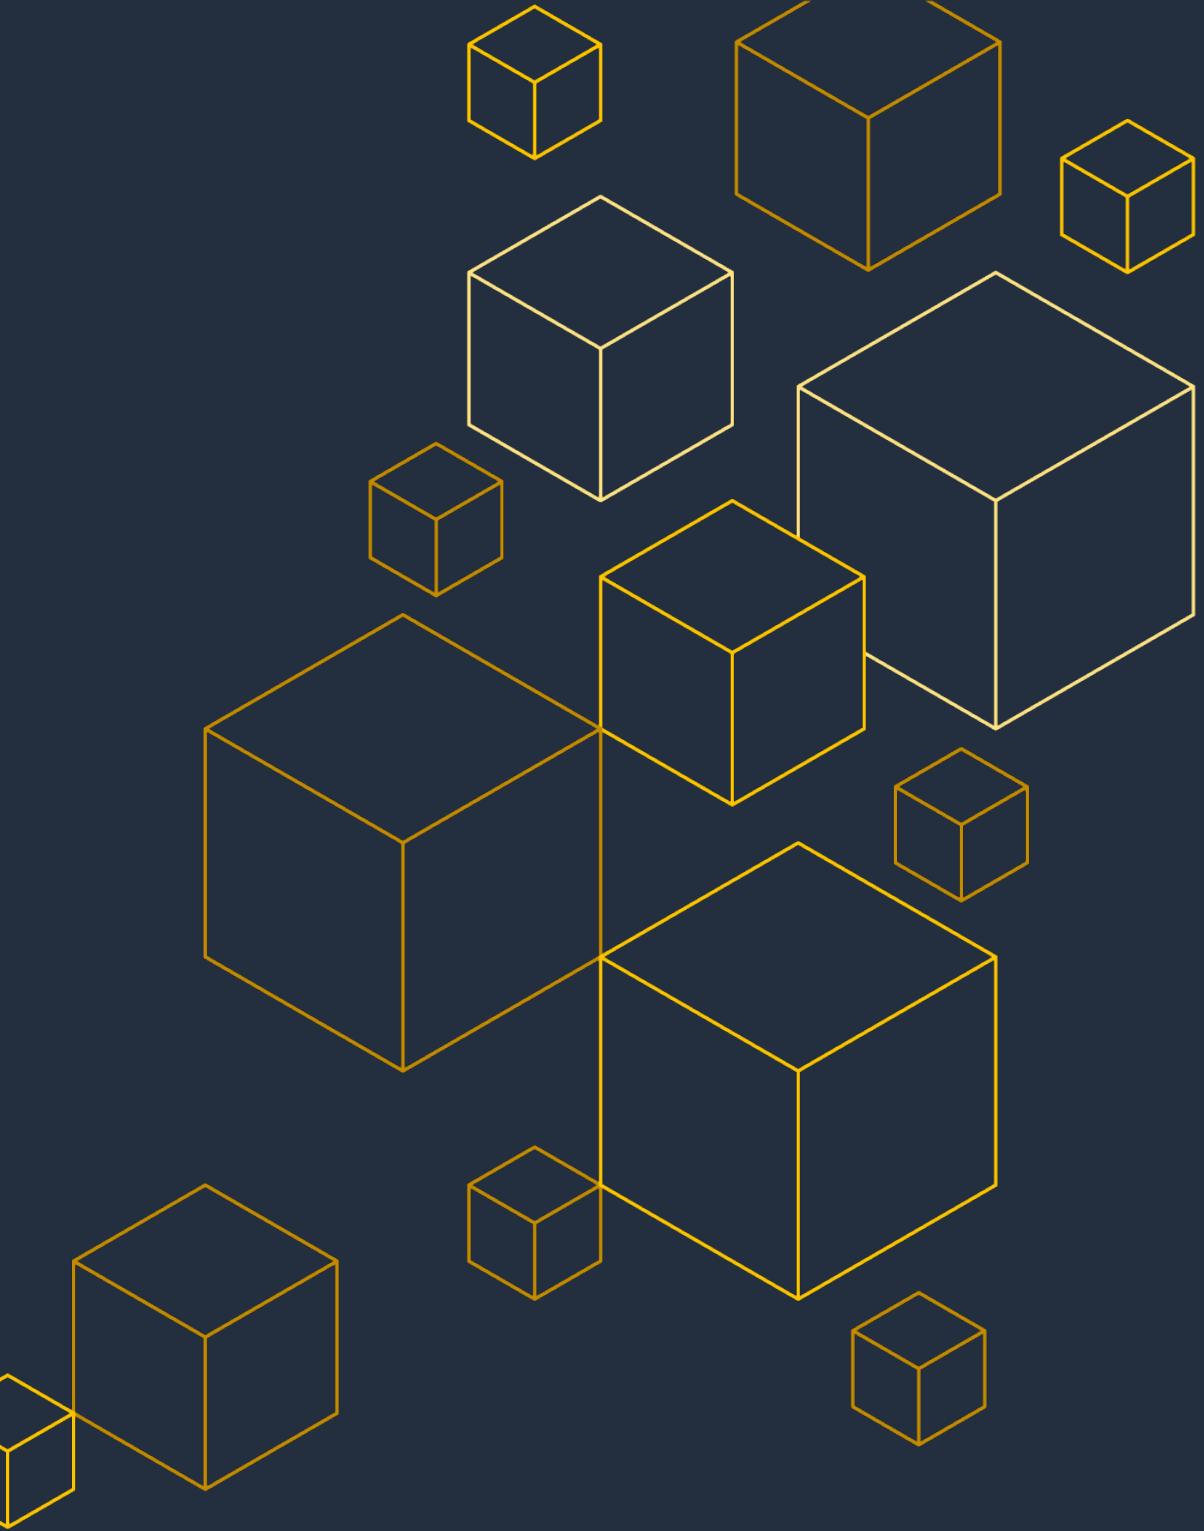
IMAGE REGISTRY

Container image repository

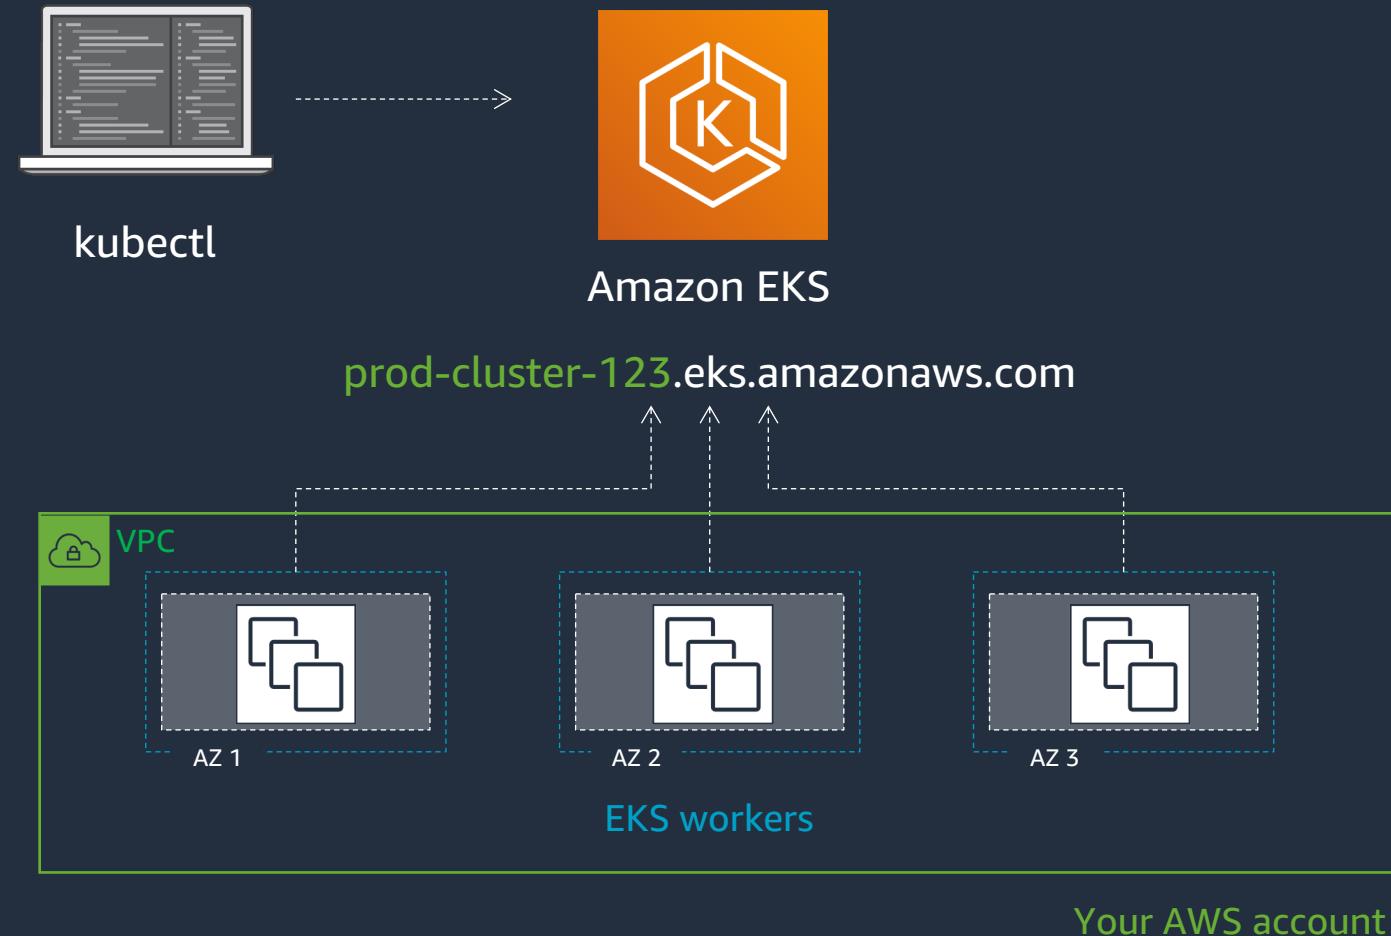


Amazon Elastic
Container Registry
(Amazon ECR)

Amazon EKS



Amazon EKS Architecture



Amazon EKS 특징



**Security
First**



**Built for
Production**



**Native and
upstream**



**Seamless
Cloud
integrations**



**Committed to
Open Source**

Amazon ECS



Amazon ECS 구성 요소



ECS 클러스터

작업이 실행되는 논리적 그룹



ECS 작업

실제 컨테이너 작업(최소 실행 단위)

하나 혹은 둘 이상의 컨테이너의 묶음

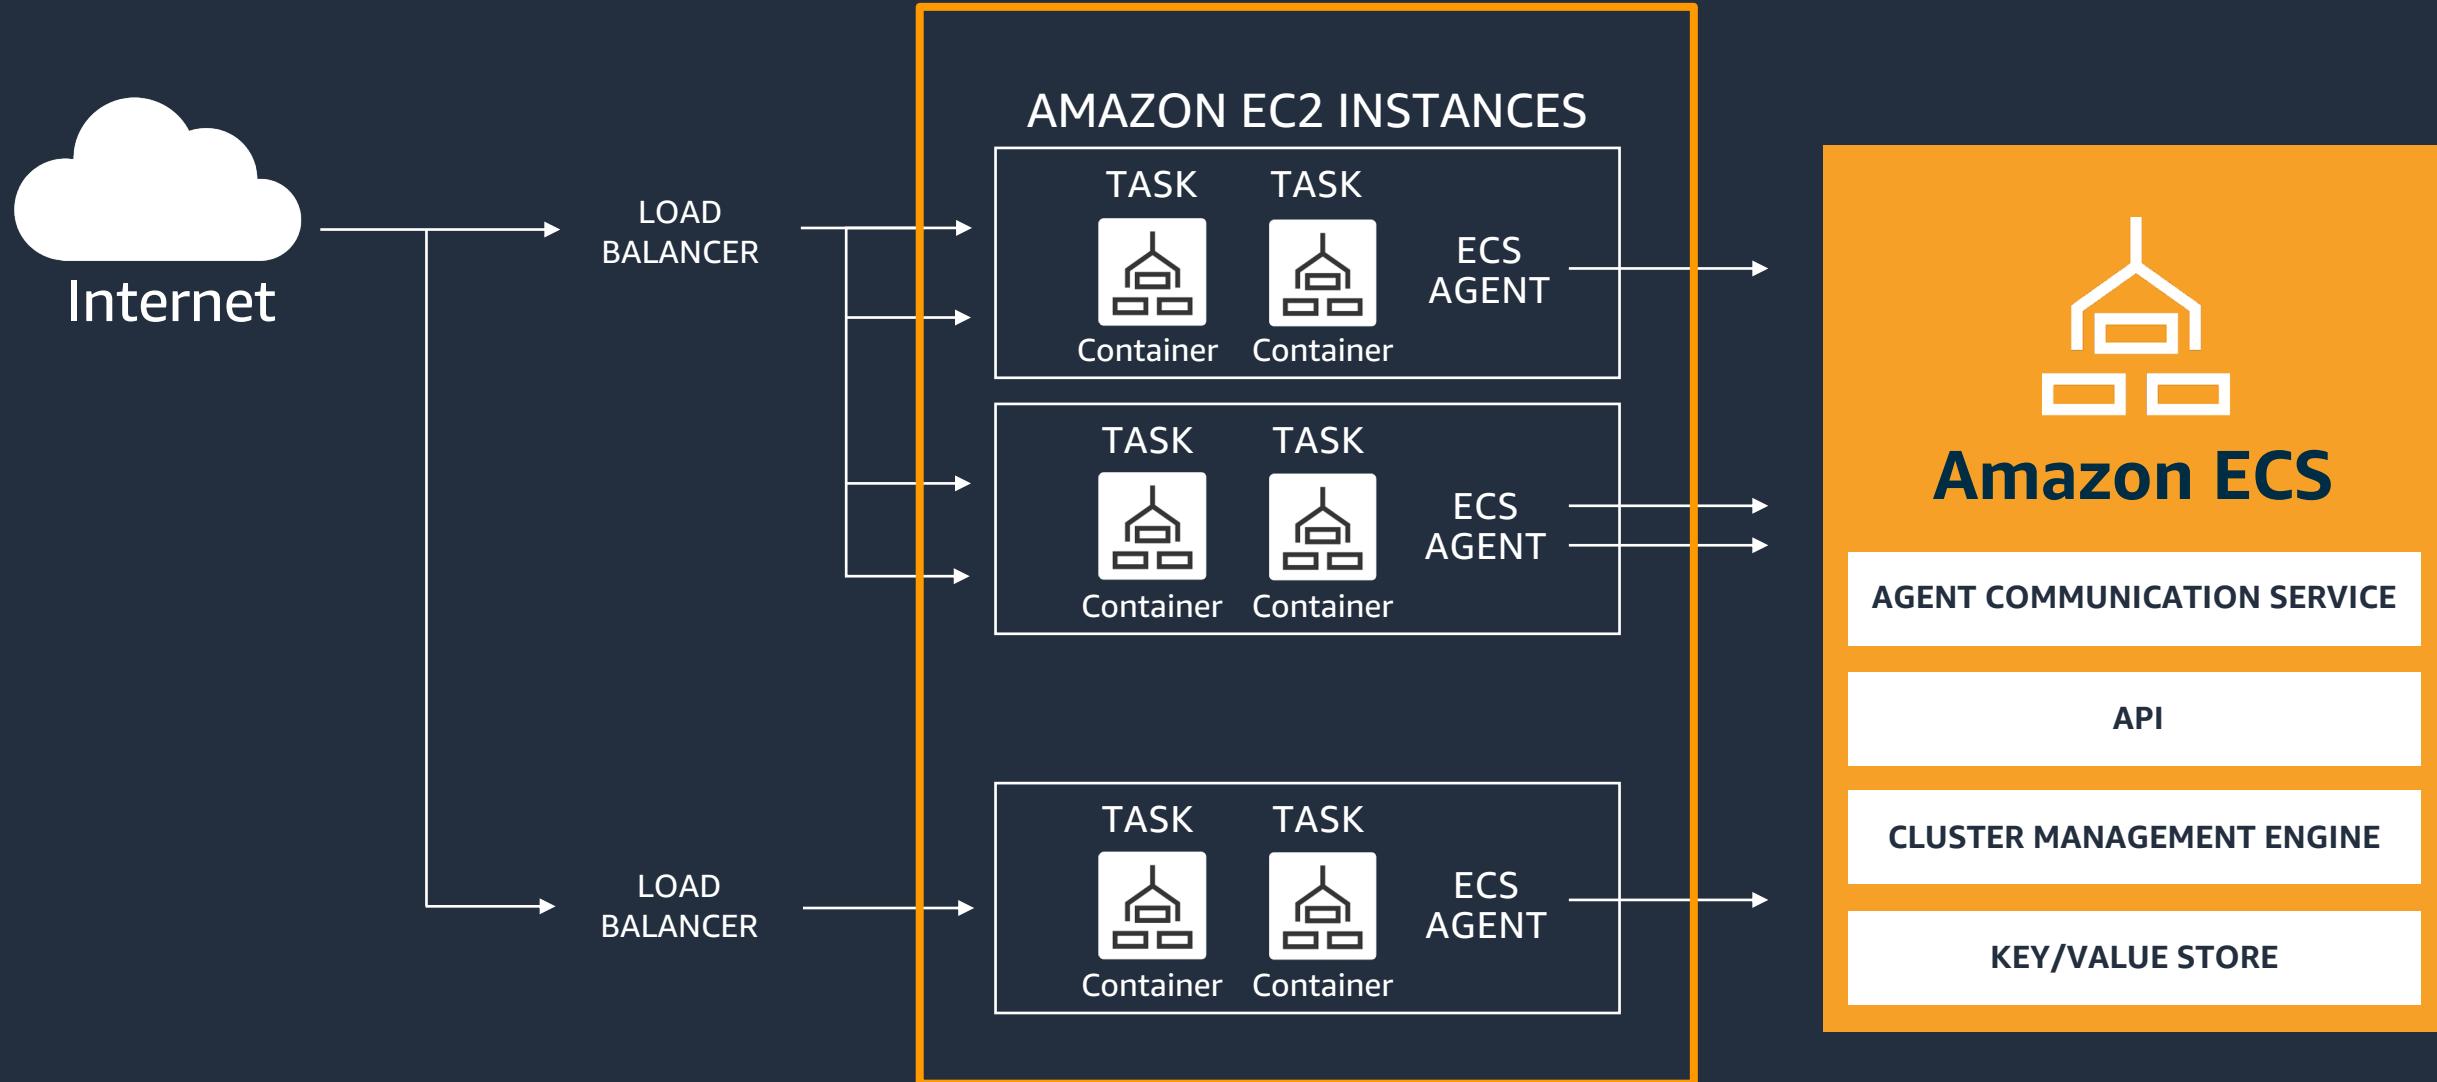


ECS 서비스

작업을 지속적으로 관리

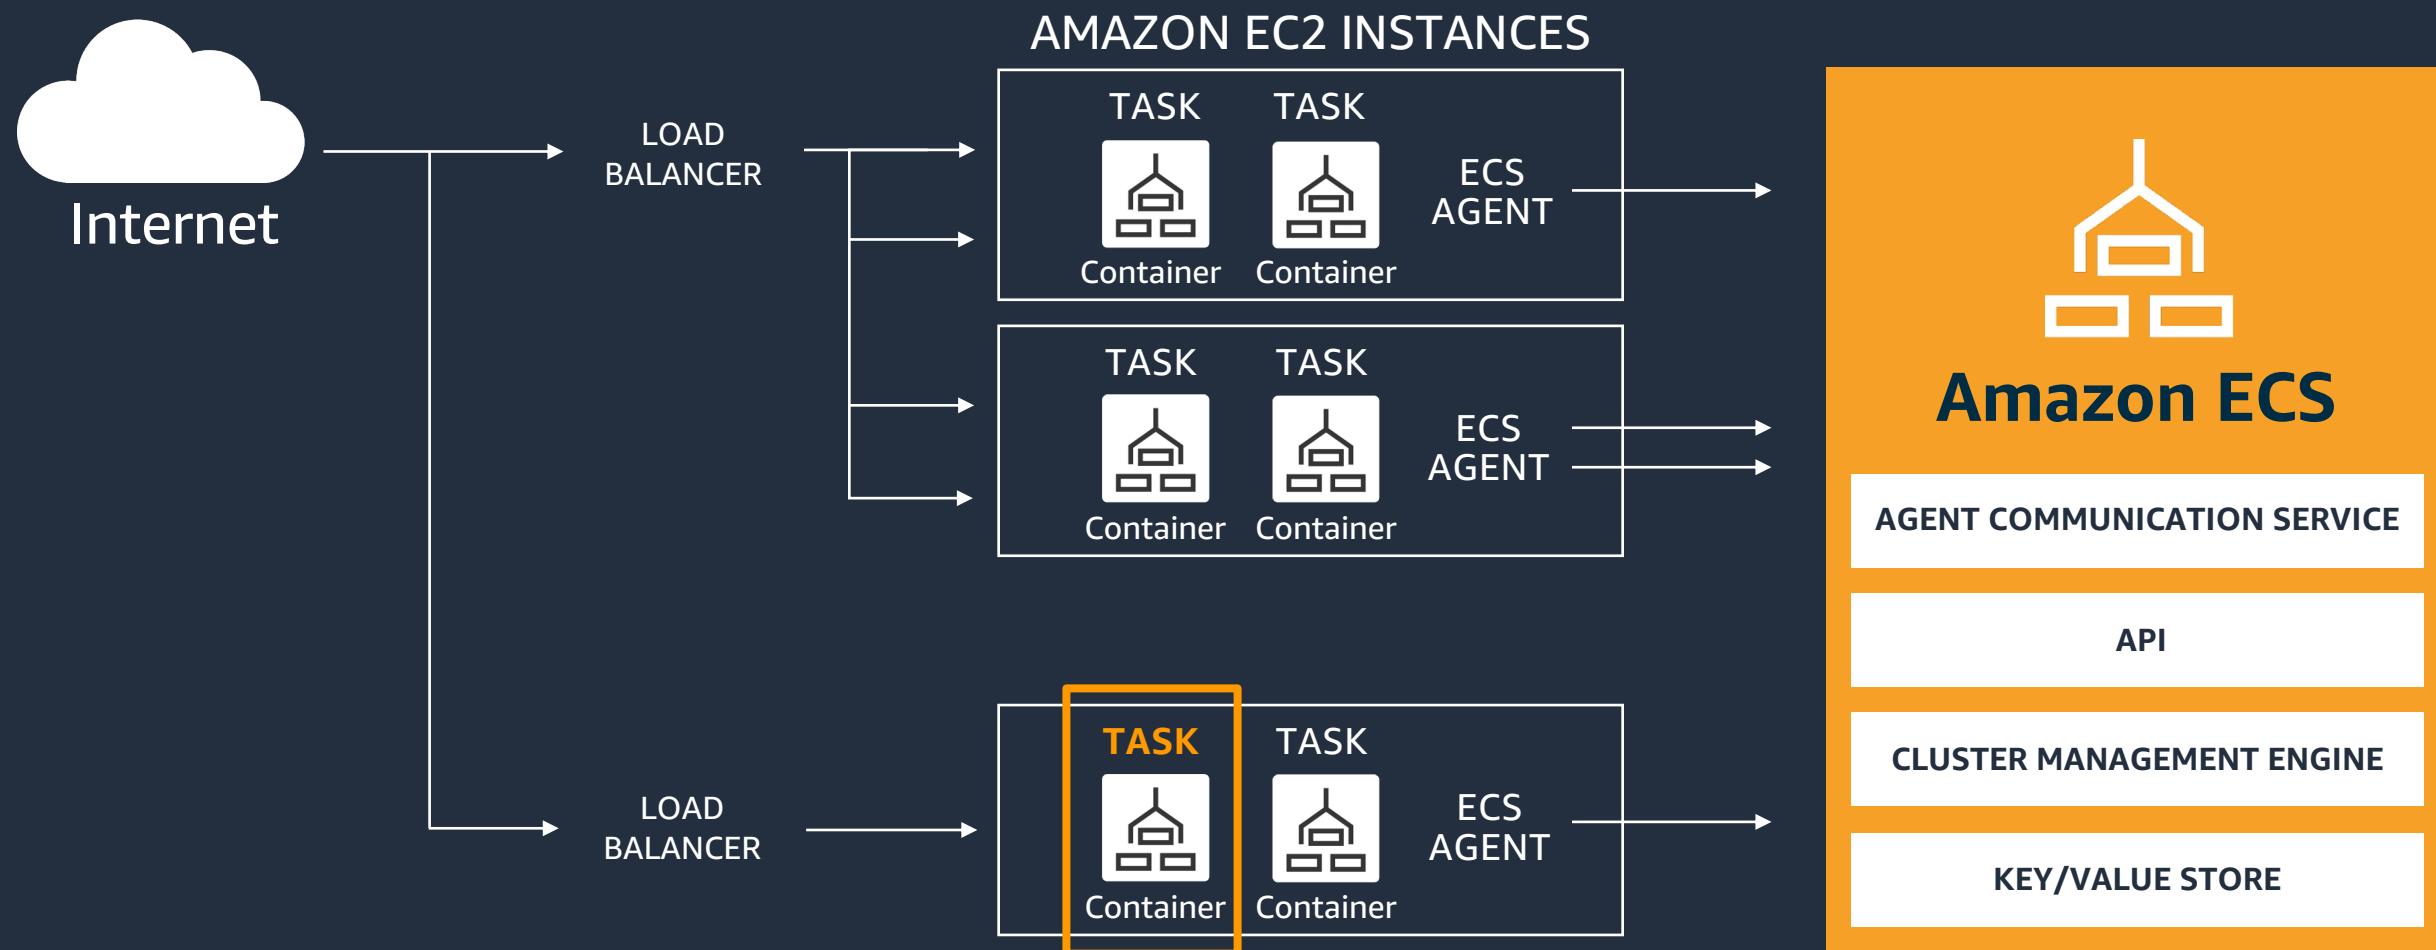
실패한 작업을 자동으로 대체(작업 수 유지)

Amazon ECS



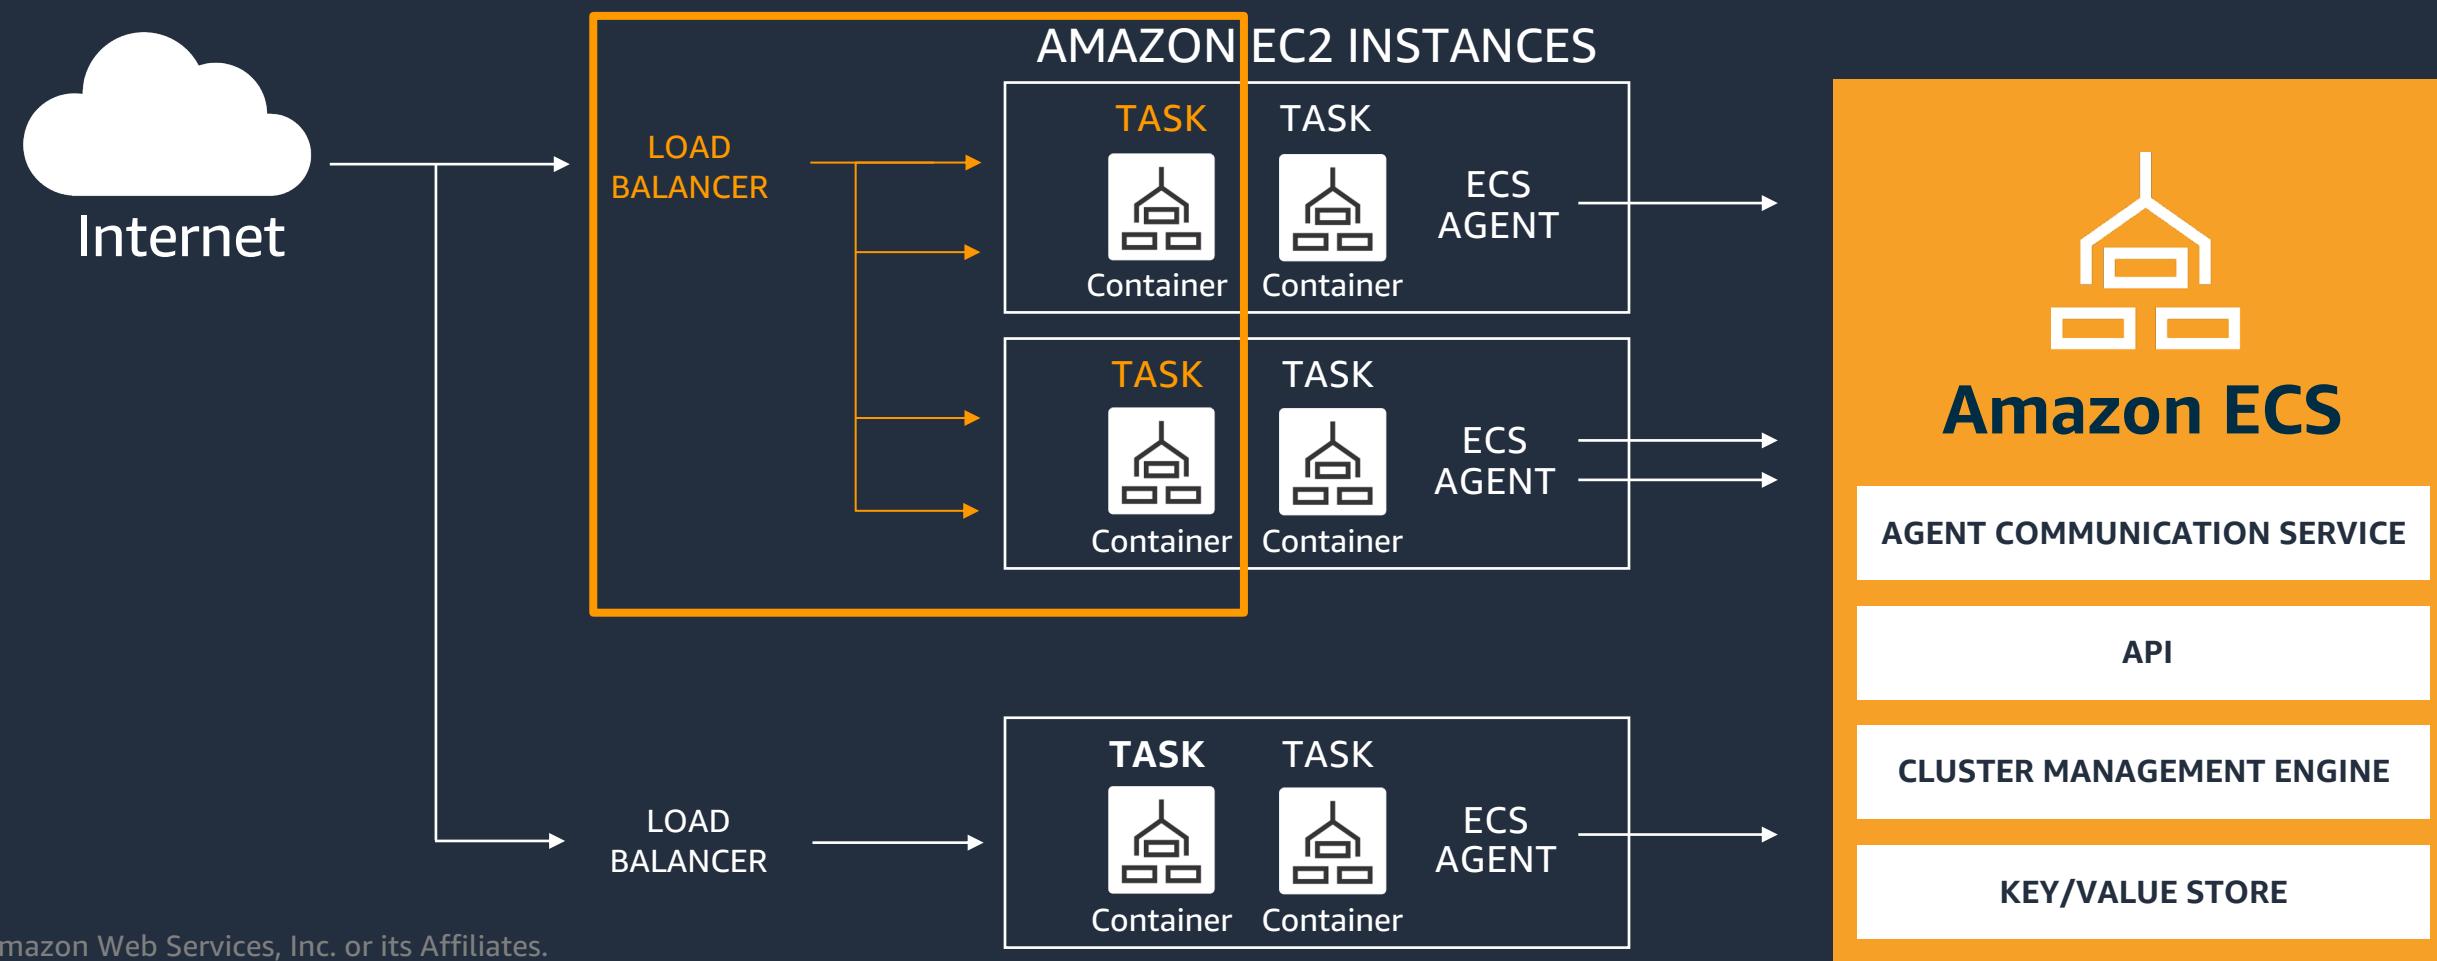
Amazon ECS - 작업

- ECS에서 **최소 실행 단위**
- 작업 정의 내용을 기반으로 ECS 클러스터에 속하는 인스턴스 및 Fargate 배포
- 한 작업 당 최대 10개 컨테이너 가능, 모두 같은 호스트에 배치



Amazon ECS - 서비스

- 설정한 작업 수를 자동 유지 및 복구
- 서비스 유형(replica, daemon), 작업 배치 전략 선택 가능
- 배포 유형: 롤링 업데이트(ECS로 제어), 블루/그린 배포(CodeDeploy로 제어)
- ALB를 통해 효율적으로 부하 분산
- 서비스 Auto Scaling 적용 가능



작업 정의 (Task Def)

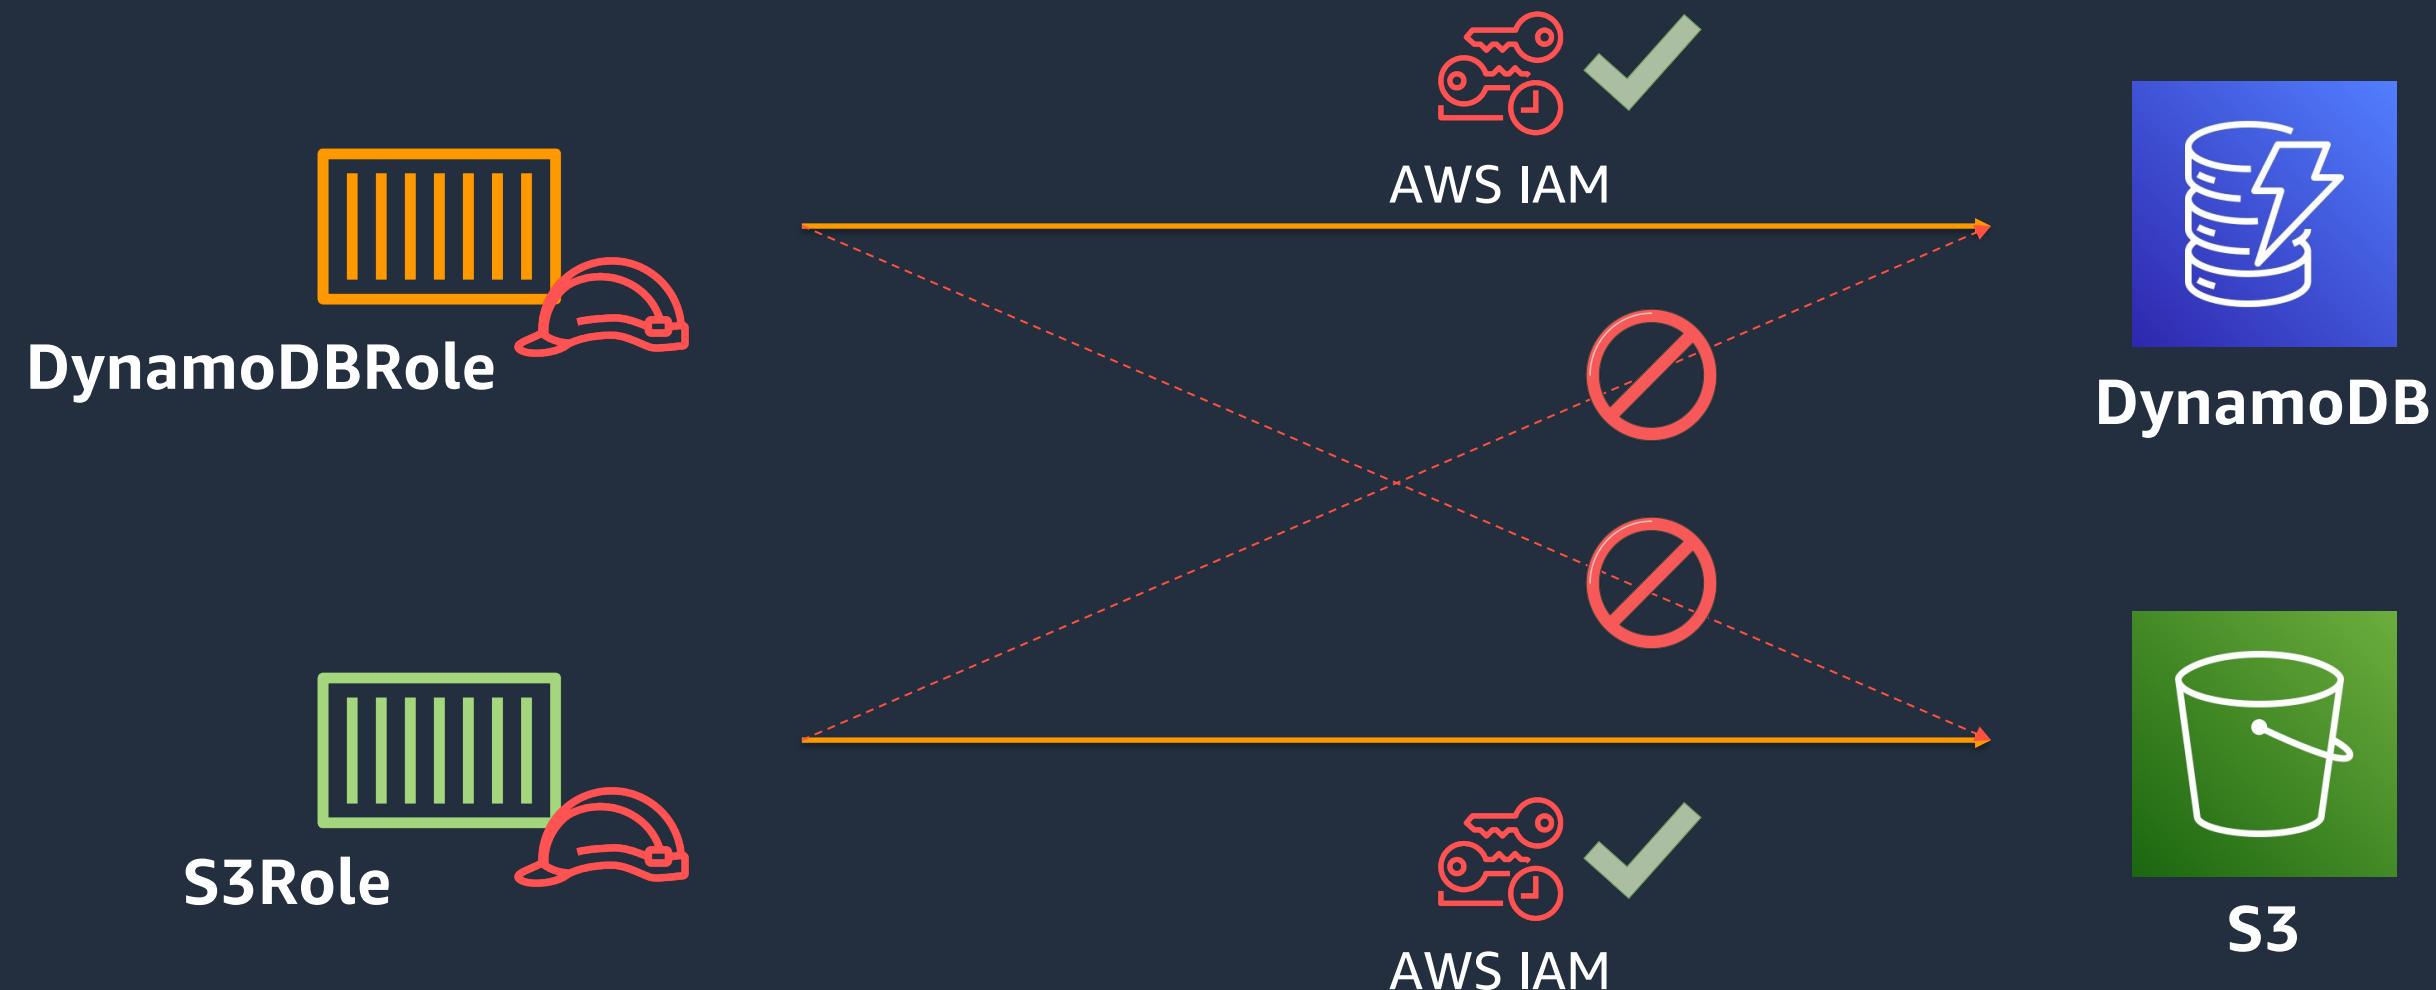


- ECS에서 컨테이너를 실행하려면 작업 정의가 필요
- 컨테이너 이미지 맵핑을 통한 컨테이너 정의
- 배포 타입 설정 가능: Fargate, EC2
- 작업 역할을 부여해 API 요청을 받을 때 권한에 따라 동작 가능
- 작업 크기 설정 가능
- App Mesh, FireLens 등과 통합 활성화 가능
- 다양한 볼륨 선택권 제공

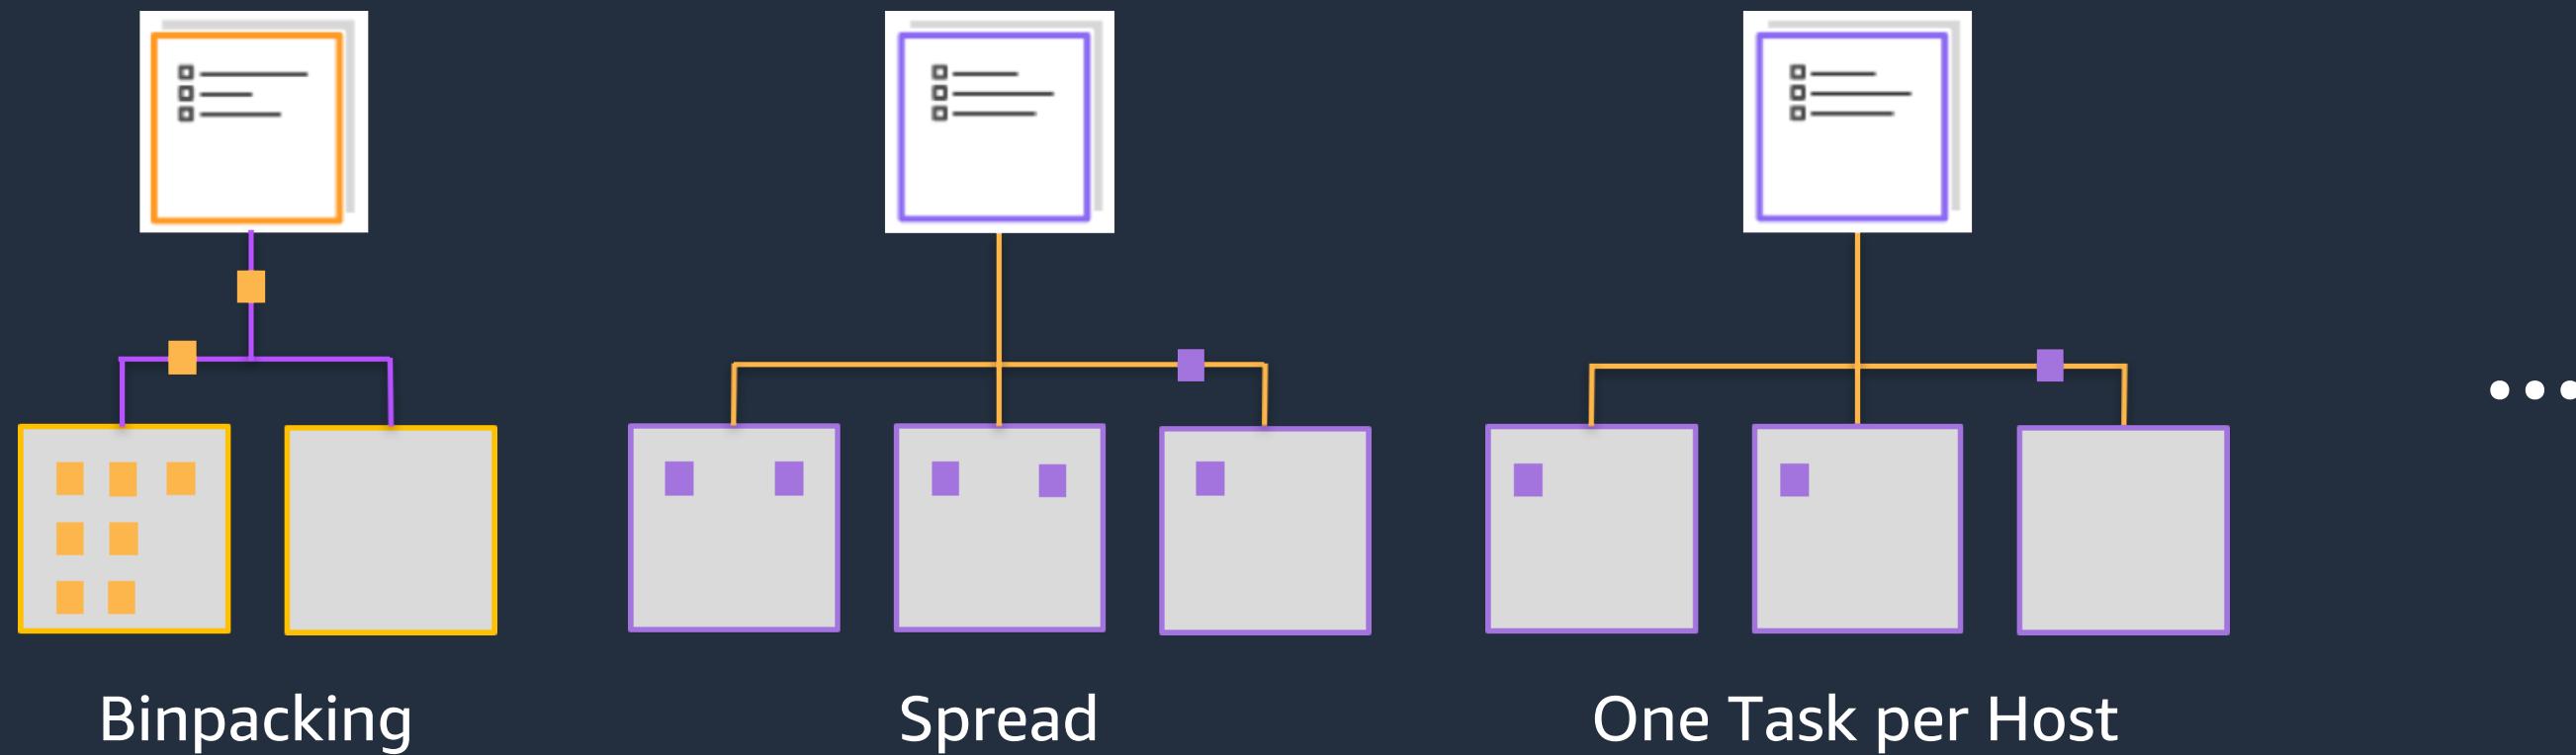
작업 정의 (Task Def)

```
{  
  "containerDefinitions": [  
    {  
      "name": "myApp",  
      "image": "myApp:v12",  
      "essential": true,  
      "portMappings": [  
        {  
          "containerPort": 80,  
          "hostPort": 80  
        }  
      ],  
      "readonlyRootFilesystem": true,  
      "memory": 500,  
      "cpu": 100,  
      ...  
      "logConfiguration": {  
        "logDriver": "awslogs",  
        "options": {  
          "awslogs-group": "ecs-log-streaming",  
          "awslogs-region": "us-west-2"  
        }  
      }  
    },  
    {  
      "family": "hello_world",  
      "taskRoleArn": "arn:aws:iam::123456789012:role/MyTaskRole"  
    }  
  ]  
}
```

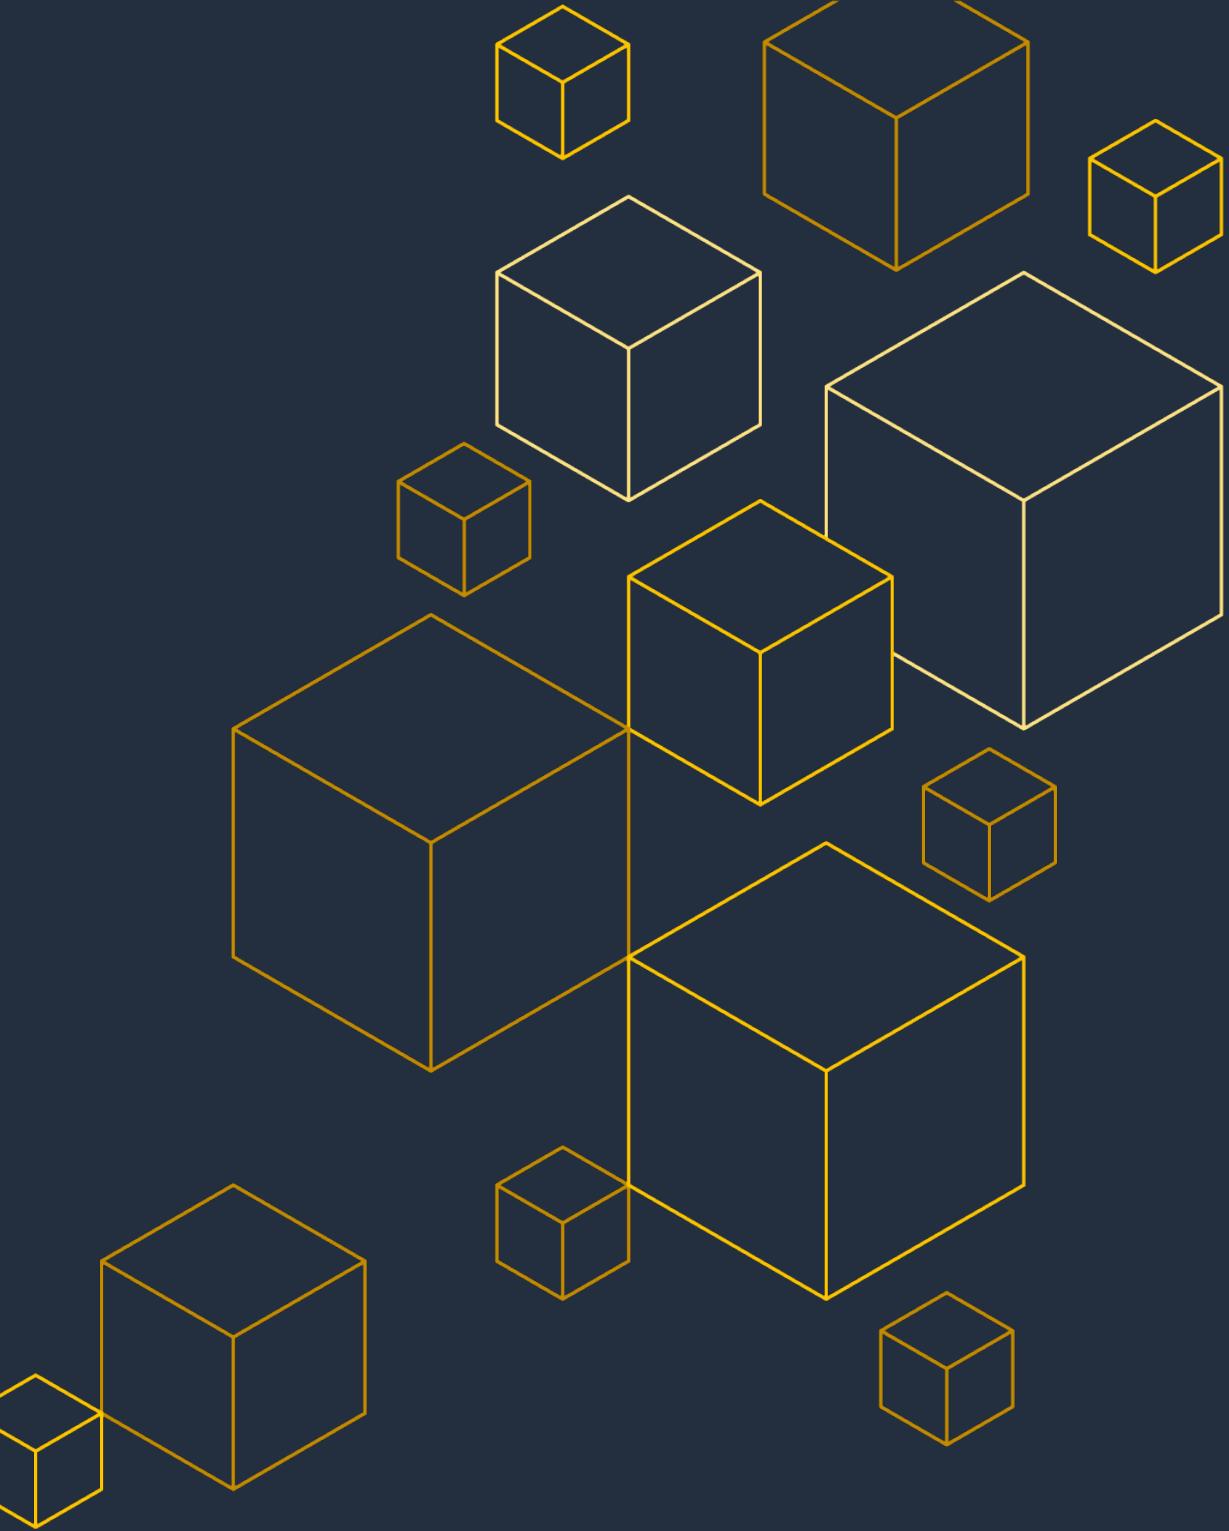
ECS 작업에 부여하는 IAM 역할



다양한 작업 배치 전략



AWS Fargate



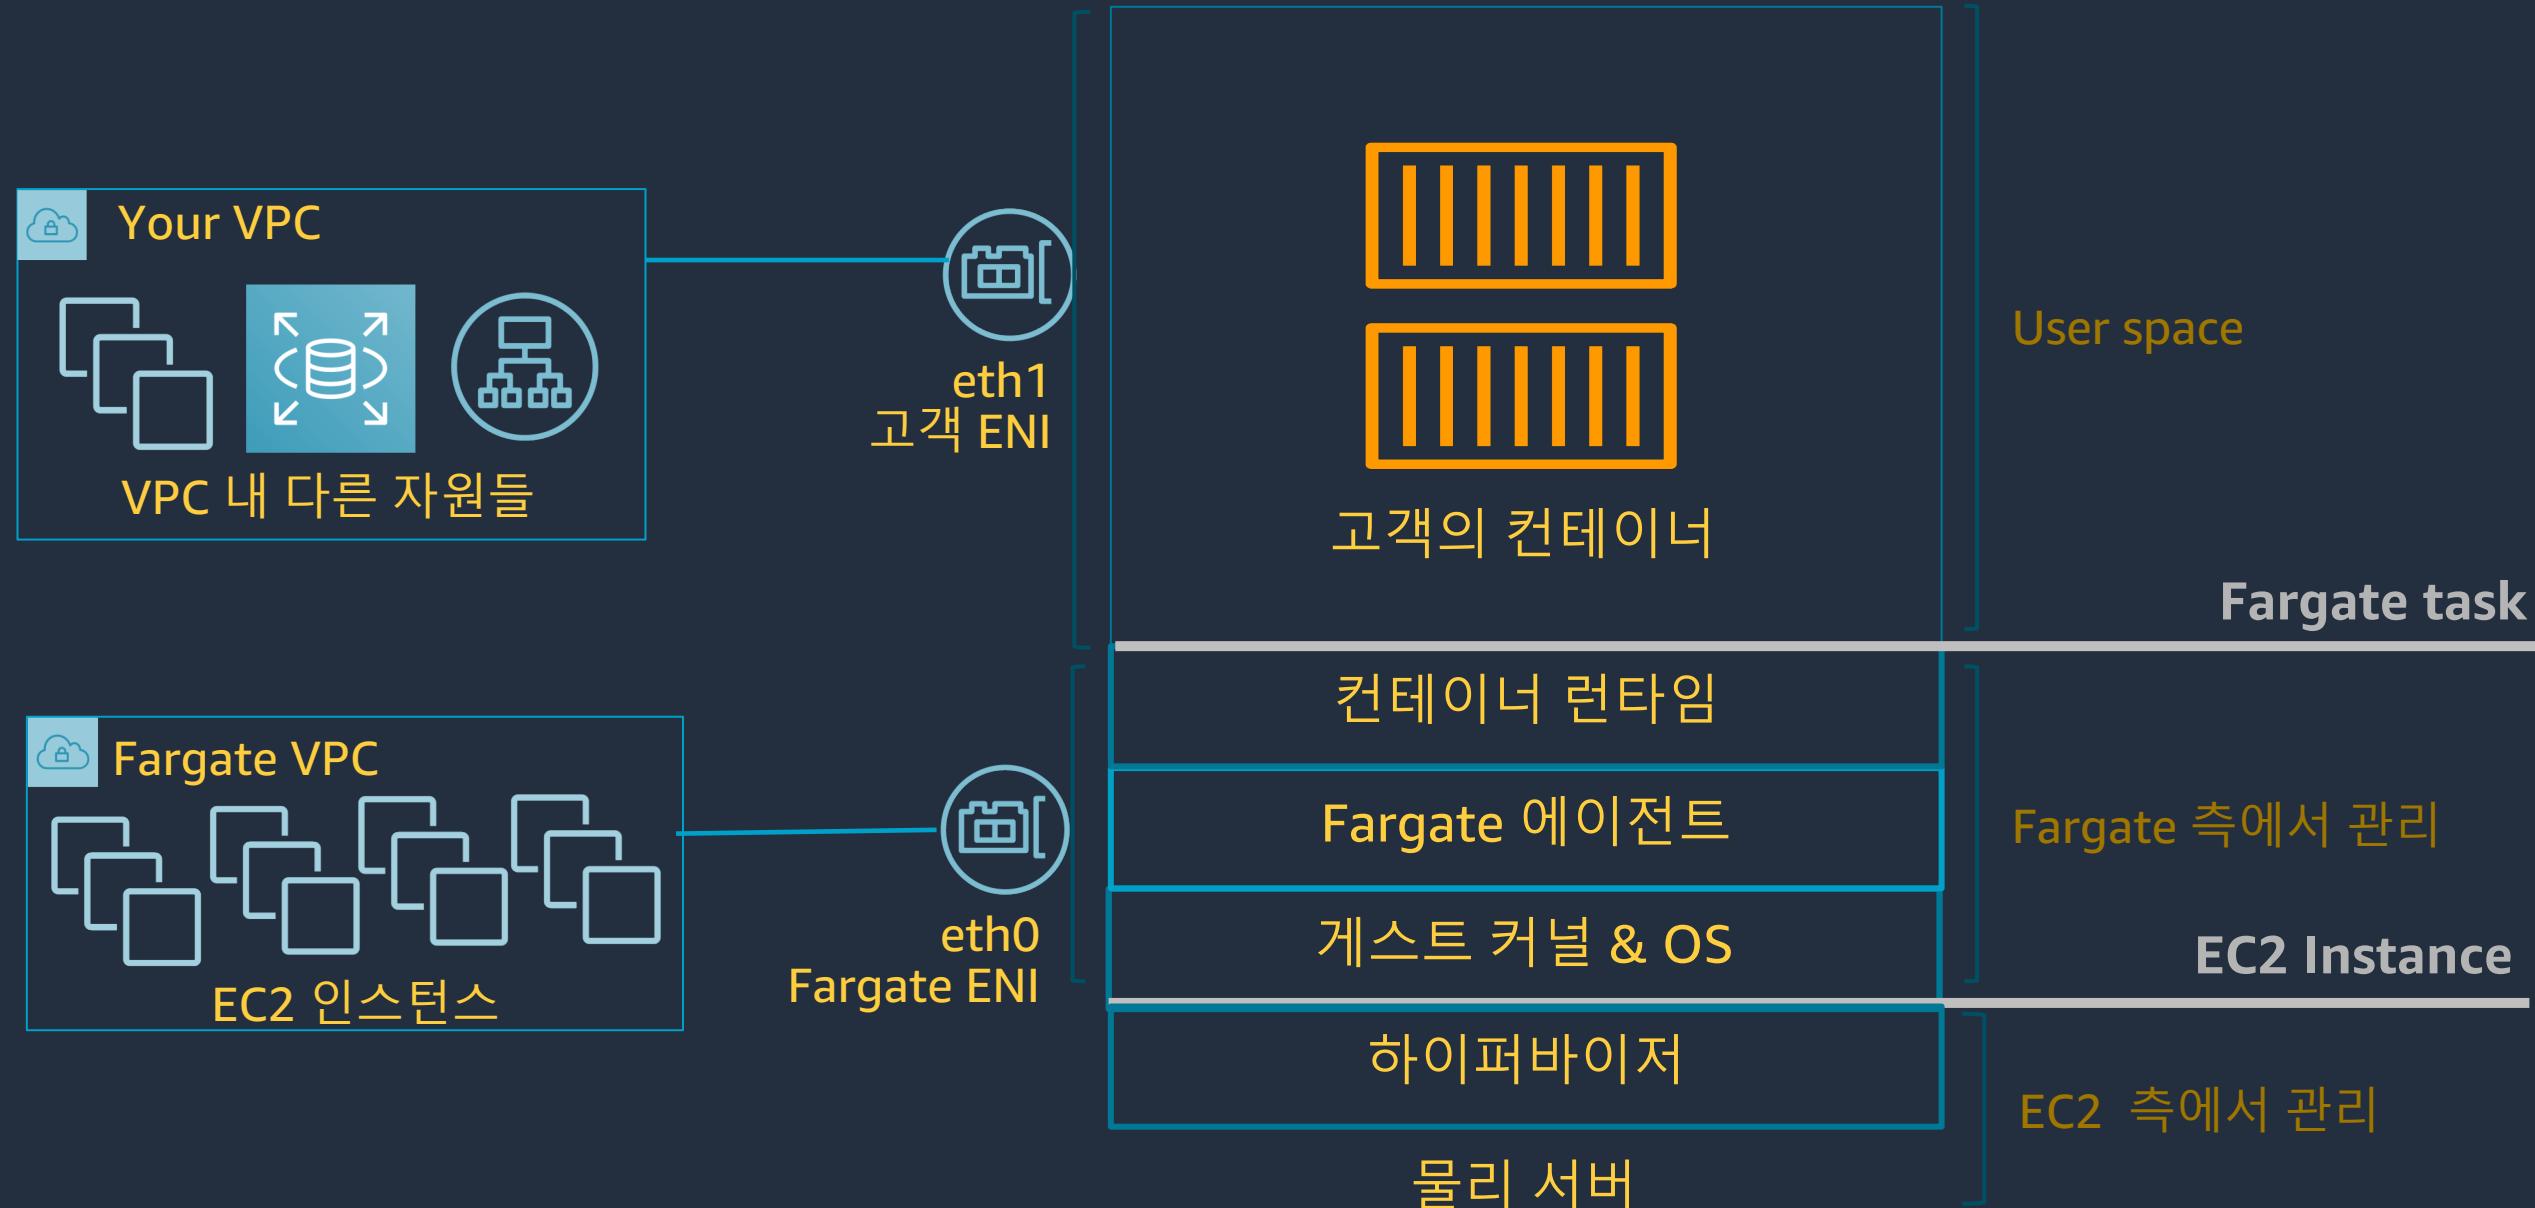
Fargate?



AWS Fargate

고객들이 기반 VM을 관리할 필요 없이
컨테이너를 사용할 수 있게 하는 컴퓨팅 엔진

Fargate Data Plane



Fargate의 장점

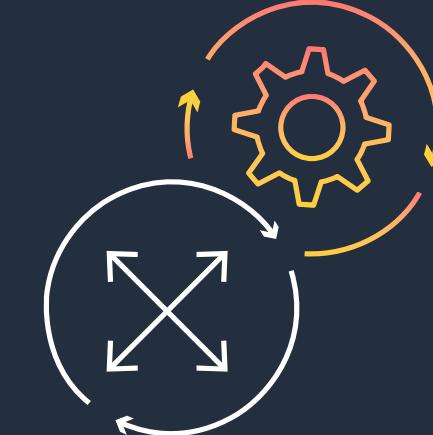


관리 업무로 인한 오버헤드 경감

까다로운 컨테이너 클러스터 관리를 AWS에 위임함으로써 고객은 애플리케이션에만 집중

기존 컨테이너 그대로 배포 가능

기존의 컨테이너 변경 불필요
현재 쿠버네티스, ECS 클러스터의 서비스, 워크플로우 그대로 이용



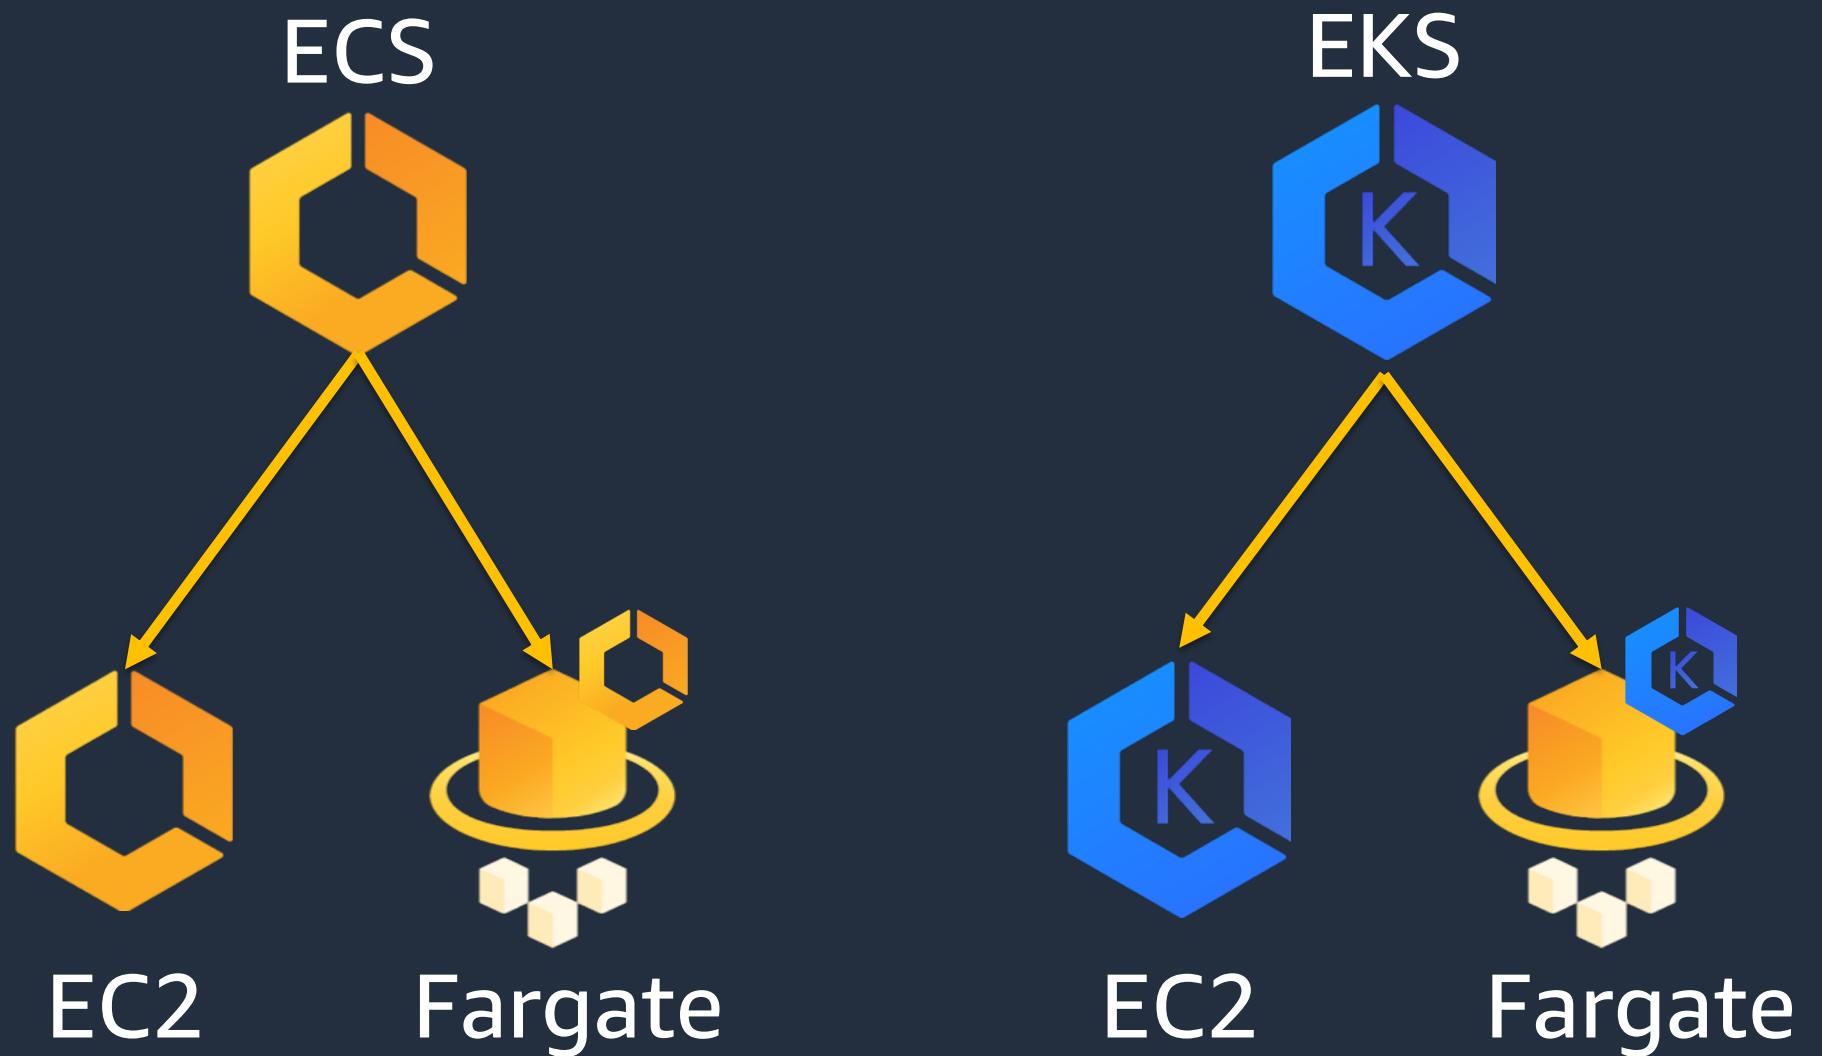
필요한 만큼만 & 쉬운 연동

컨테이너 실행에 필요한 자원 만큼만 비용 지불
기존 AWS 네트워크, 보안과 네이티브하게 통합하여 사용

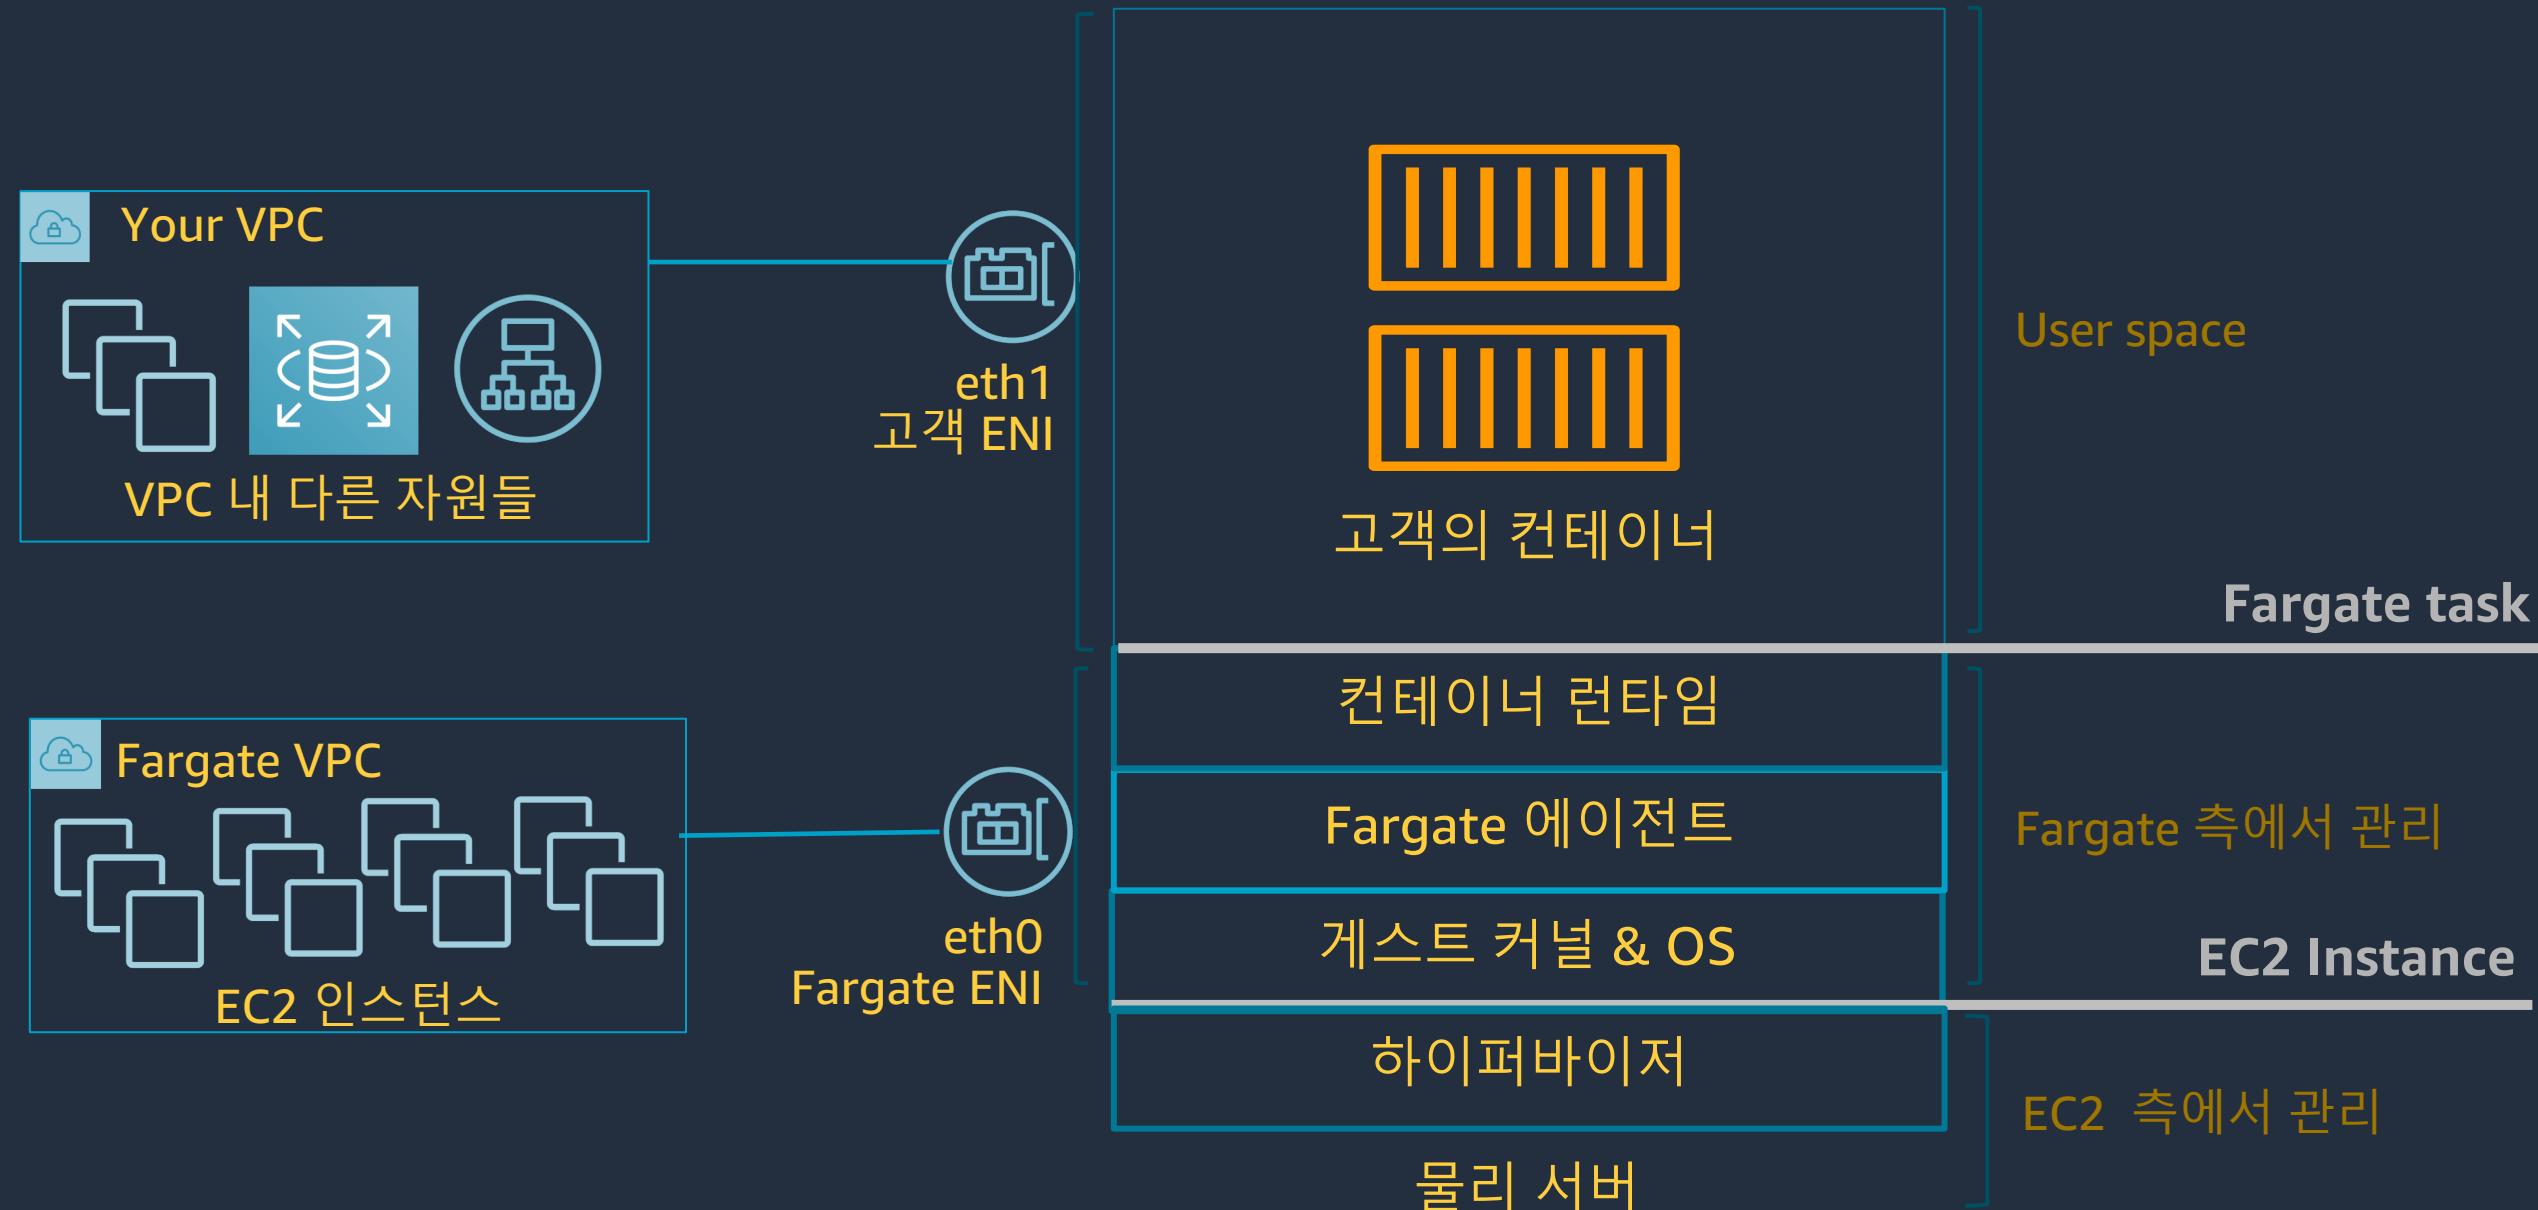
필요에 따른 조합

오케스트레이션

컴퓨트 엔진



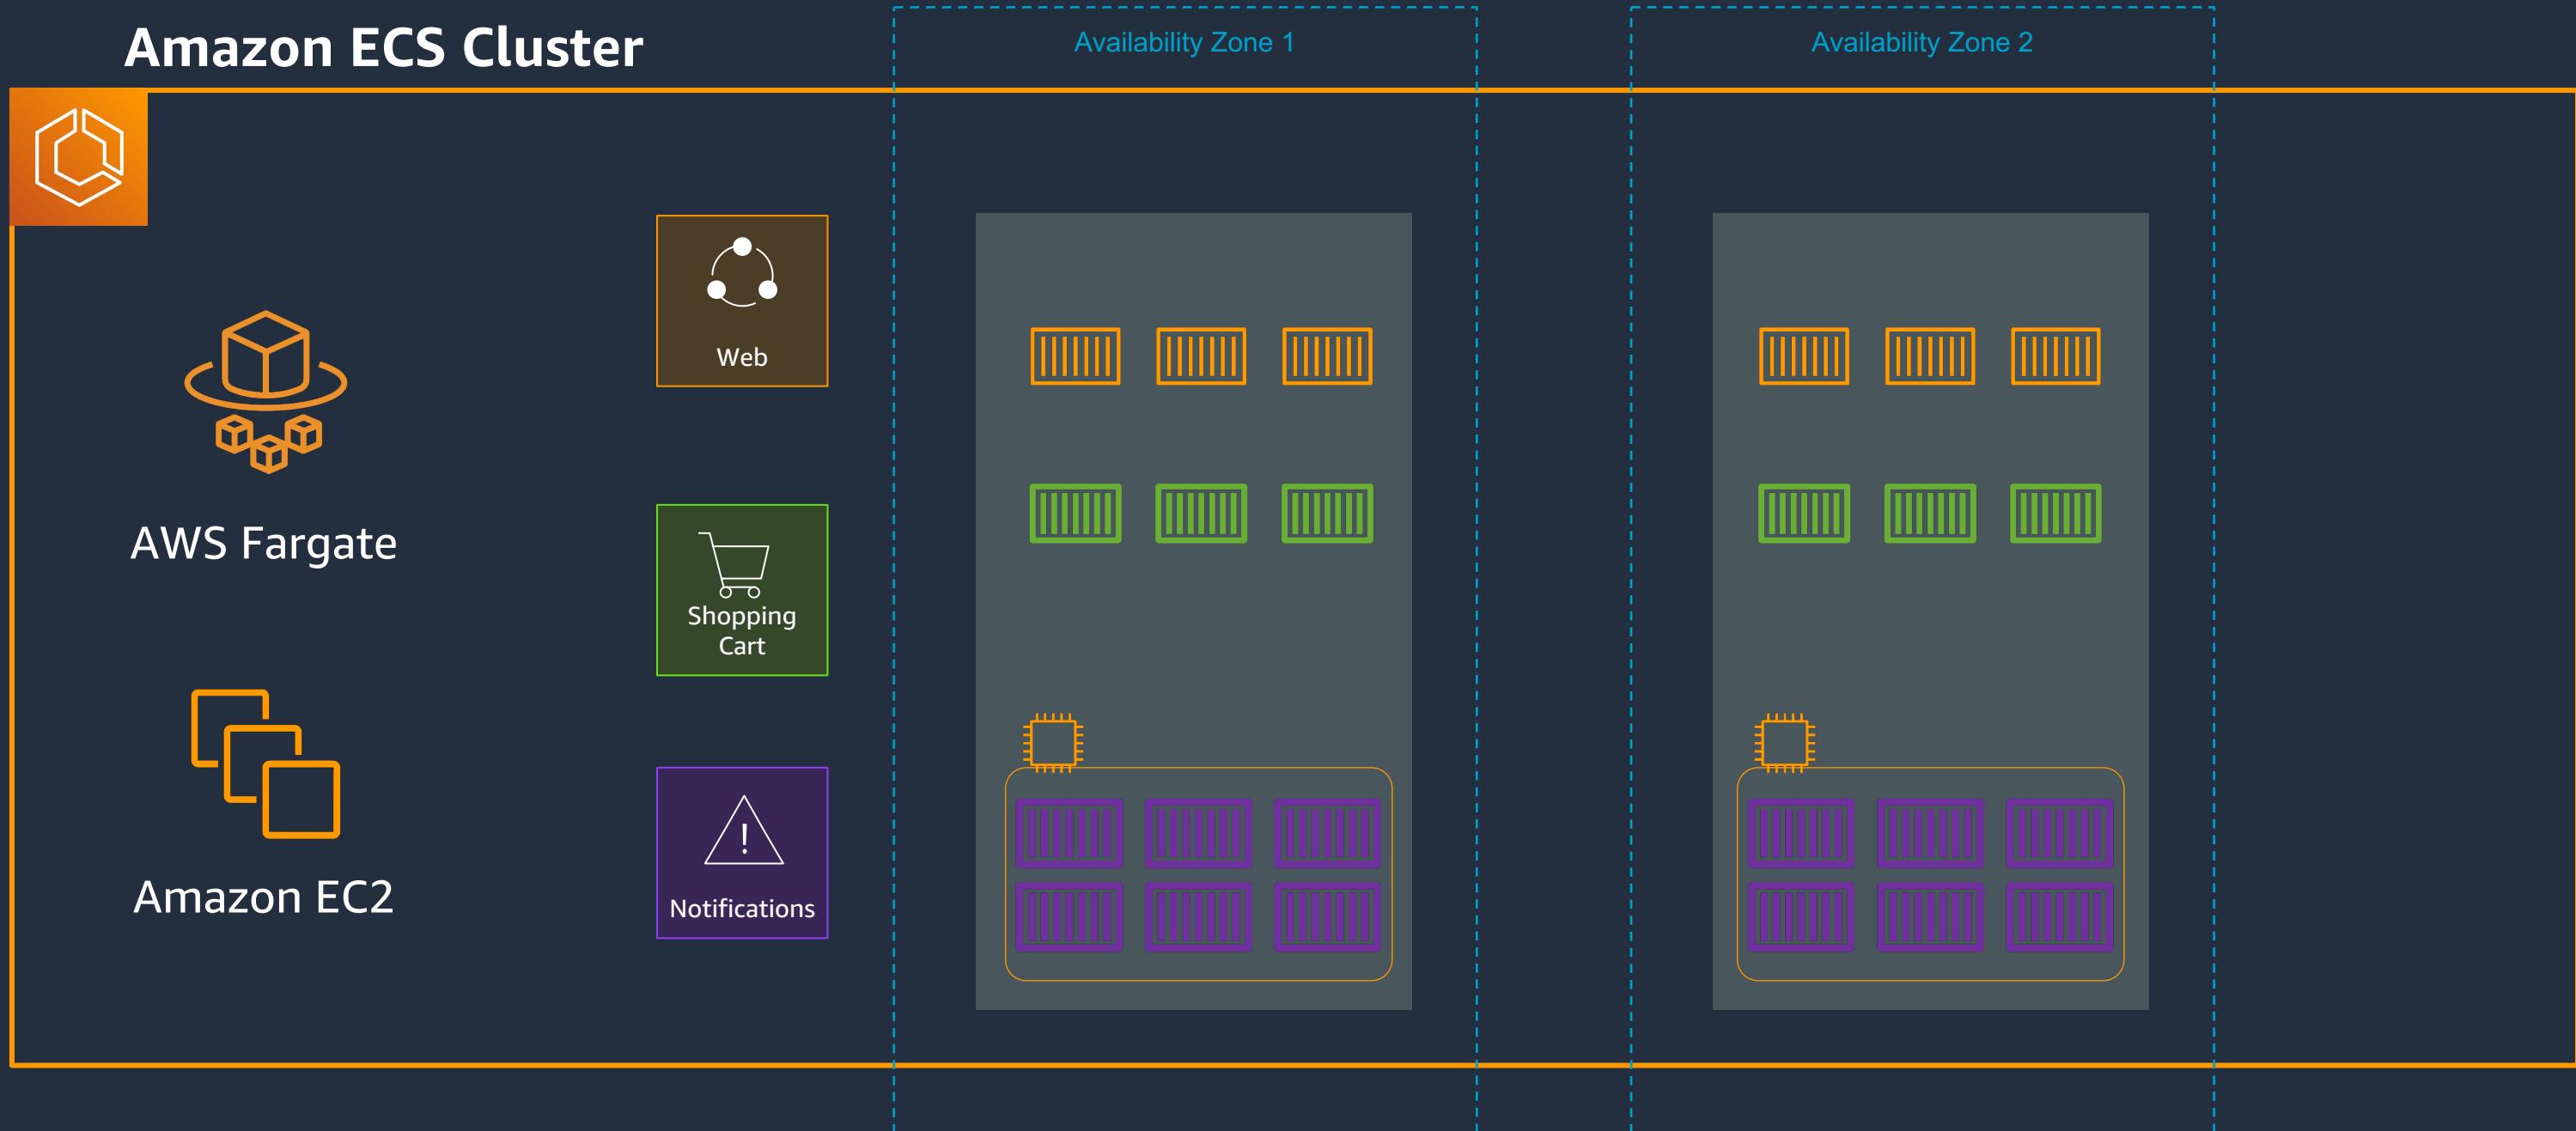
Fargate Data Plane



CPU와 메모리 용량 범위

| CPU | Memory |
|----------------|-----------------|
| 256 (.25 vCPU) | 512MB, 1GB, 2GB |
| 512 (.5 vCPU) | 1GB to 4GB |
| 1024 (1 vCPU) | 2GB to 8GB |
| 2048 (2 vCPU) | 4GB to 16GB |
| 4096 (4 vCPU) | 8GB to 30GB |

Amazon EC2와 Fargate의 Hybrid Architecture



AWS 컨테이너 서비스들



Amazon Elastic Container Registry (ECR)



- 완전 관리형 컨테이너 레지스트리
- 배포 워크플로우 간소화
 - Amazon EKS, Amazon ECS, Amazon Lambda, AWS Fargate
- 권한 제어 및 수명 주기 정책 규칙 설정
- 이미지 취약성 스캐닝 기능
- 퍼블릭/프라이빗 레파지토리 및 퍼블릭 갤러리
- 교차 리전/교차 계정 복제

<https://ecs-cats-dogs.workshop.aws/ko/ecr.html>



Amazon CloudWatch

Amazon CloudWatch를 사용하여 Amazon ECS 리소스를 모니터링

Amazon CloudWatch는 Amazon ECS에서 데이터를 수집하고 처리하여 실시간에 가까운 메트릭으로 처리



**Amazon CloudWatch
Alarms**



**Amazon CloudWatch
Logs**



**Amazon CloudWatch
Events**

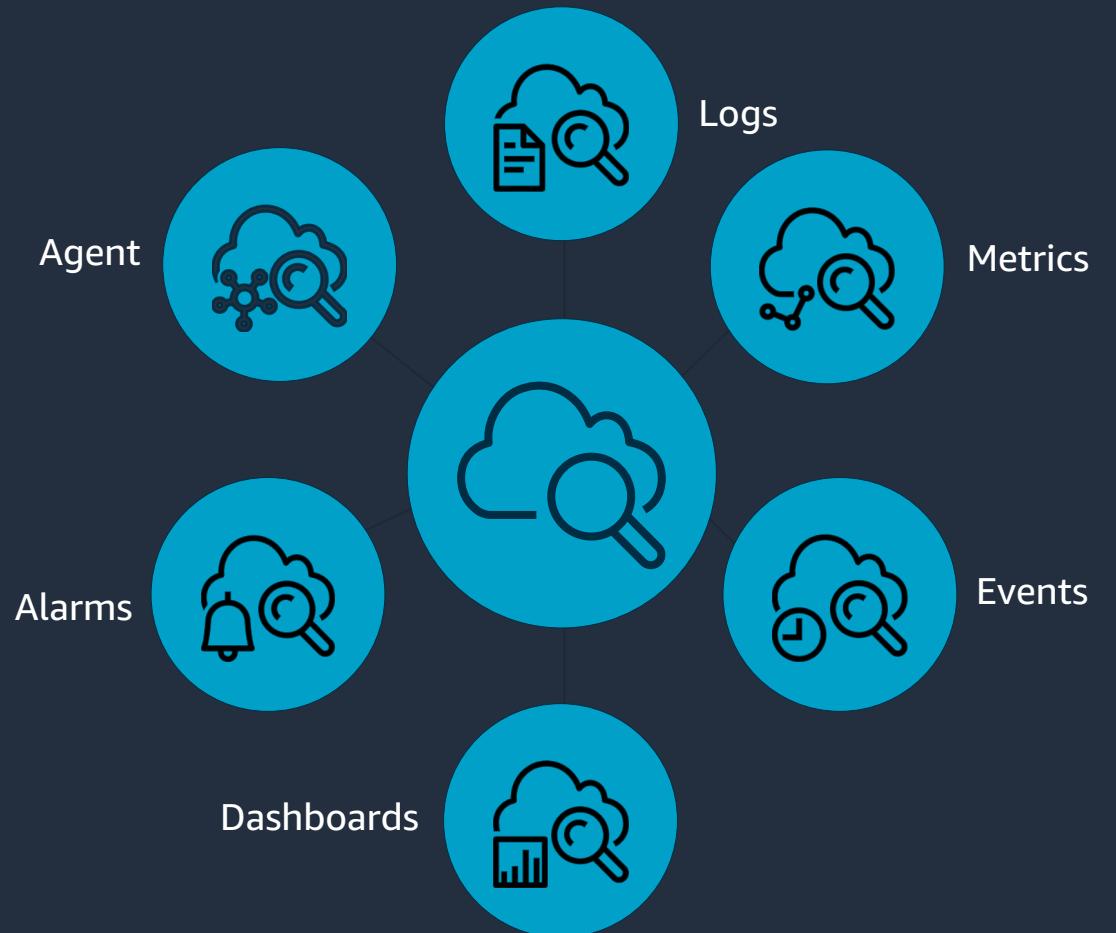


**Amazon CloudTrail
Logs**

Amazon CloudWatch Container Insights

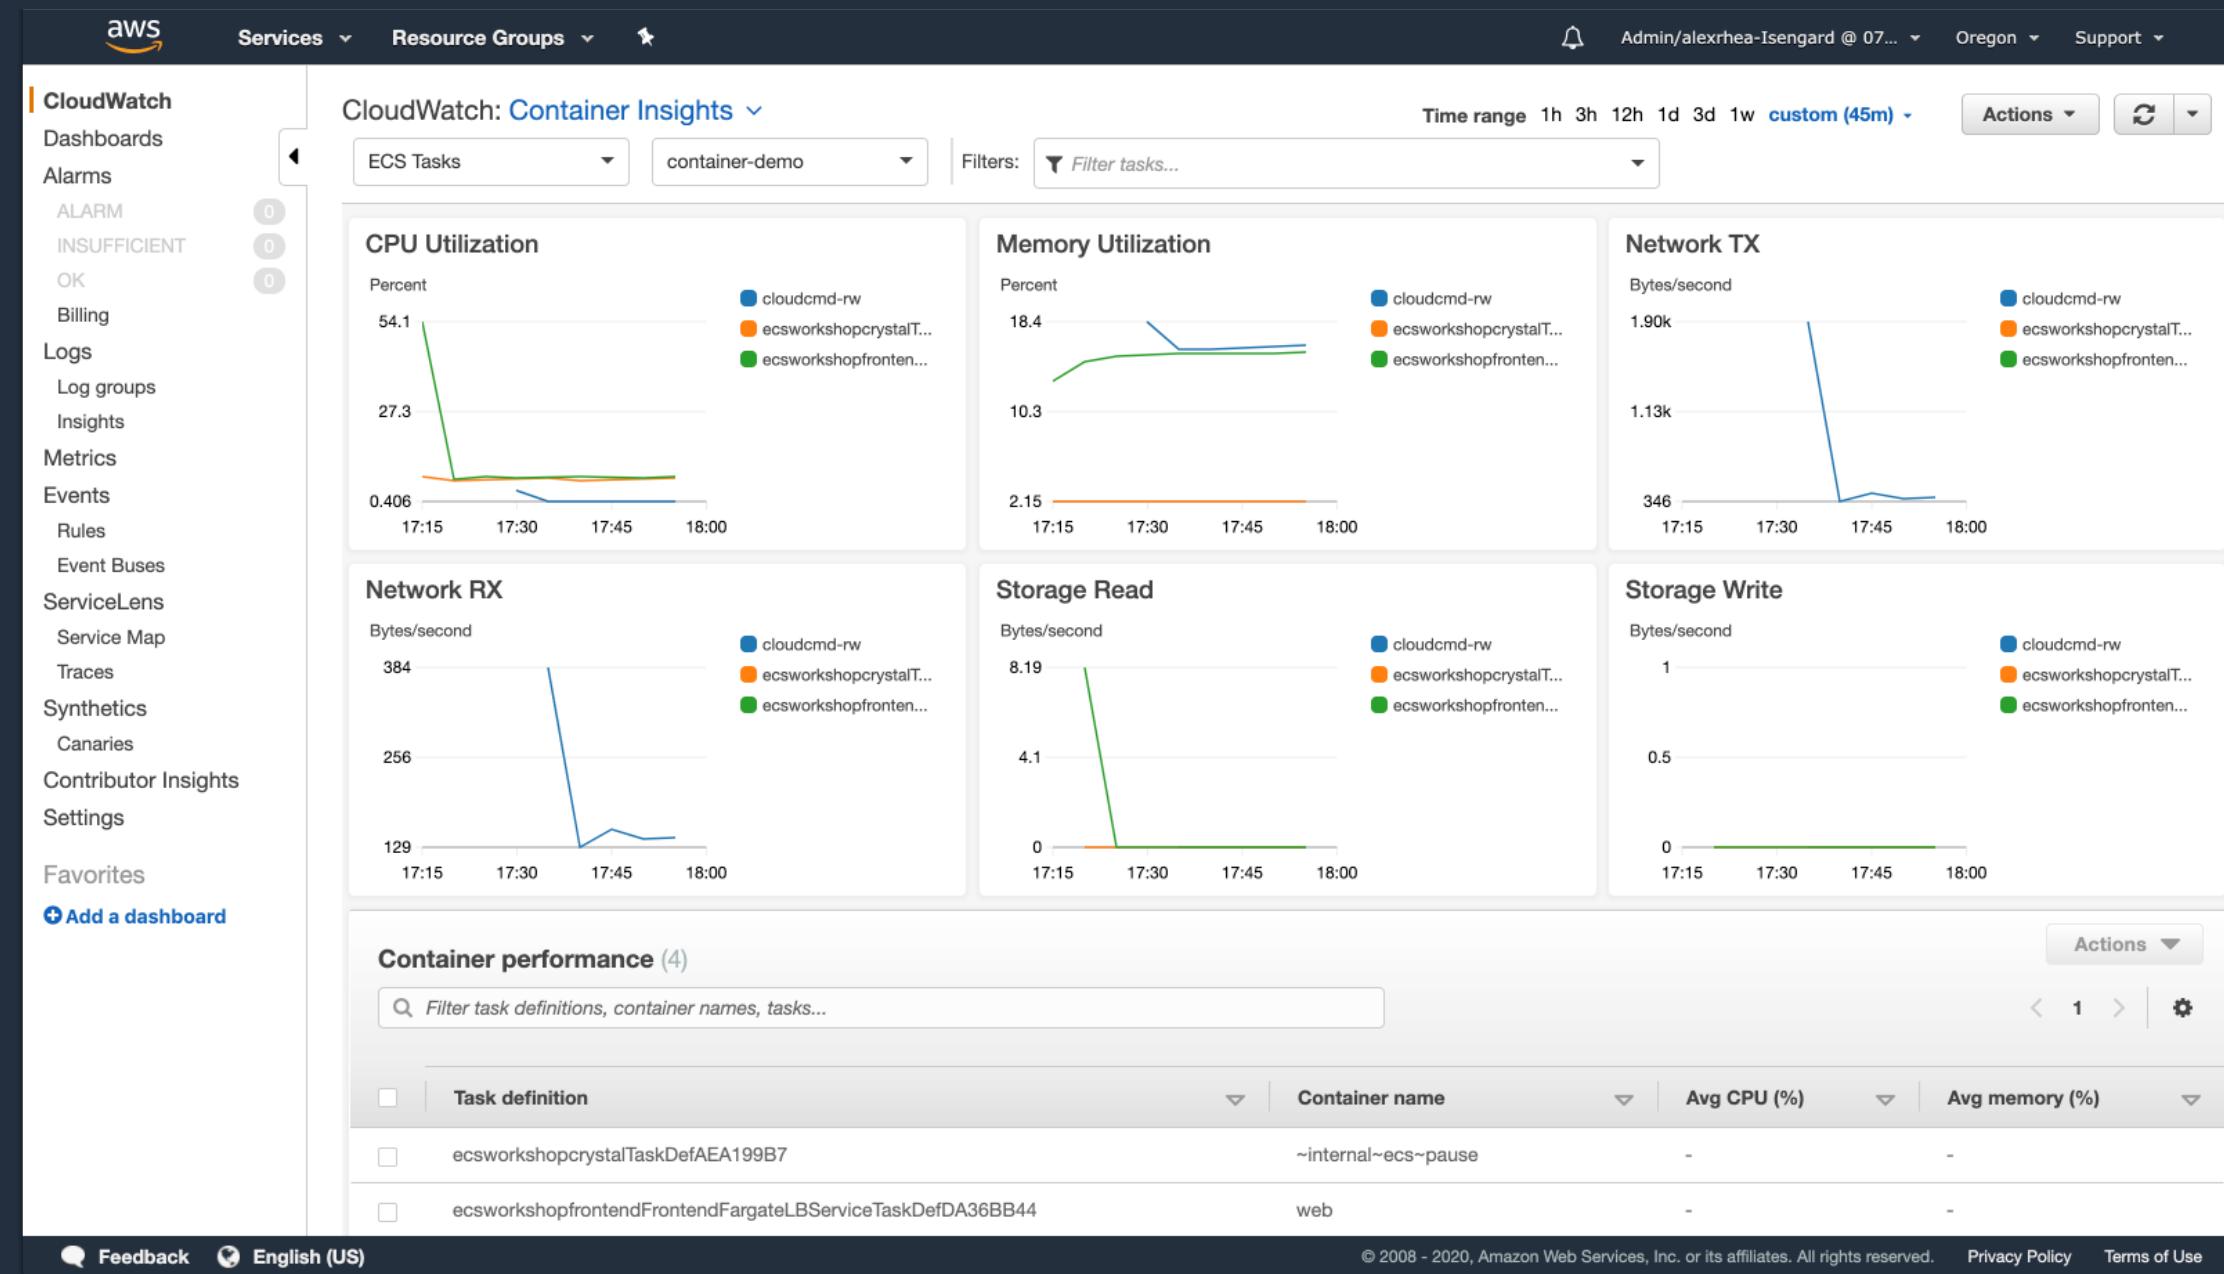
컨테이너화된 애플리케이션이나 마이크로서비스에
대한 모니터링, 트러블슈팅 및 알람을 위한 **완전**
관리형 관측 서비스

- 안정적이고 안전한 지표 및 로그 수집
- 자동화된 요약과 분석
- 지표, 로그, 트레이스 전반에 대한 가시성
- 사전 생성된 대시보드



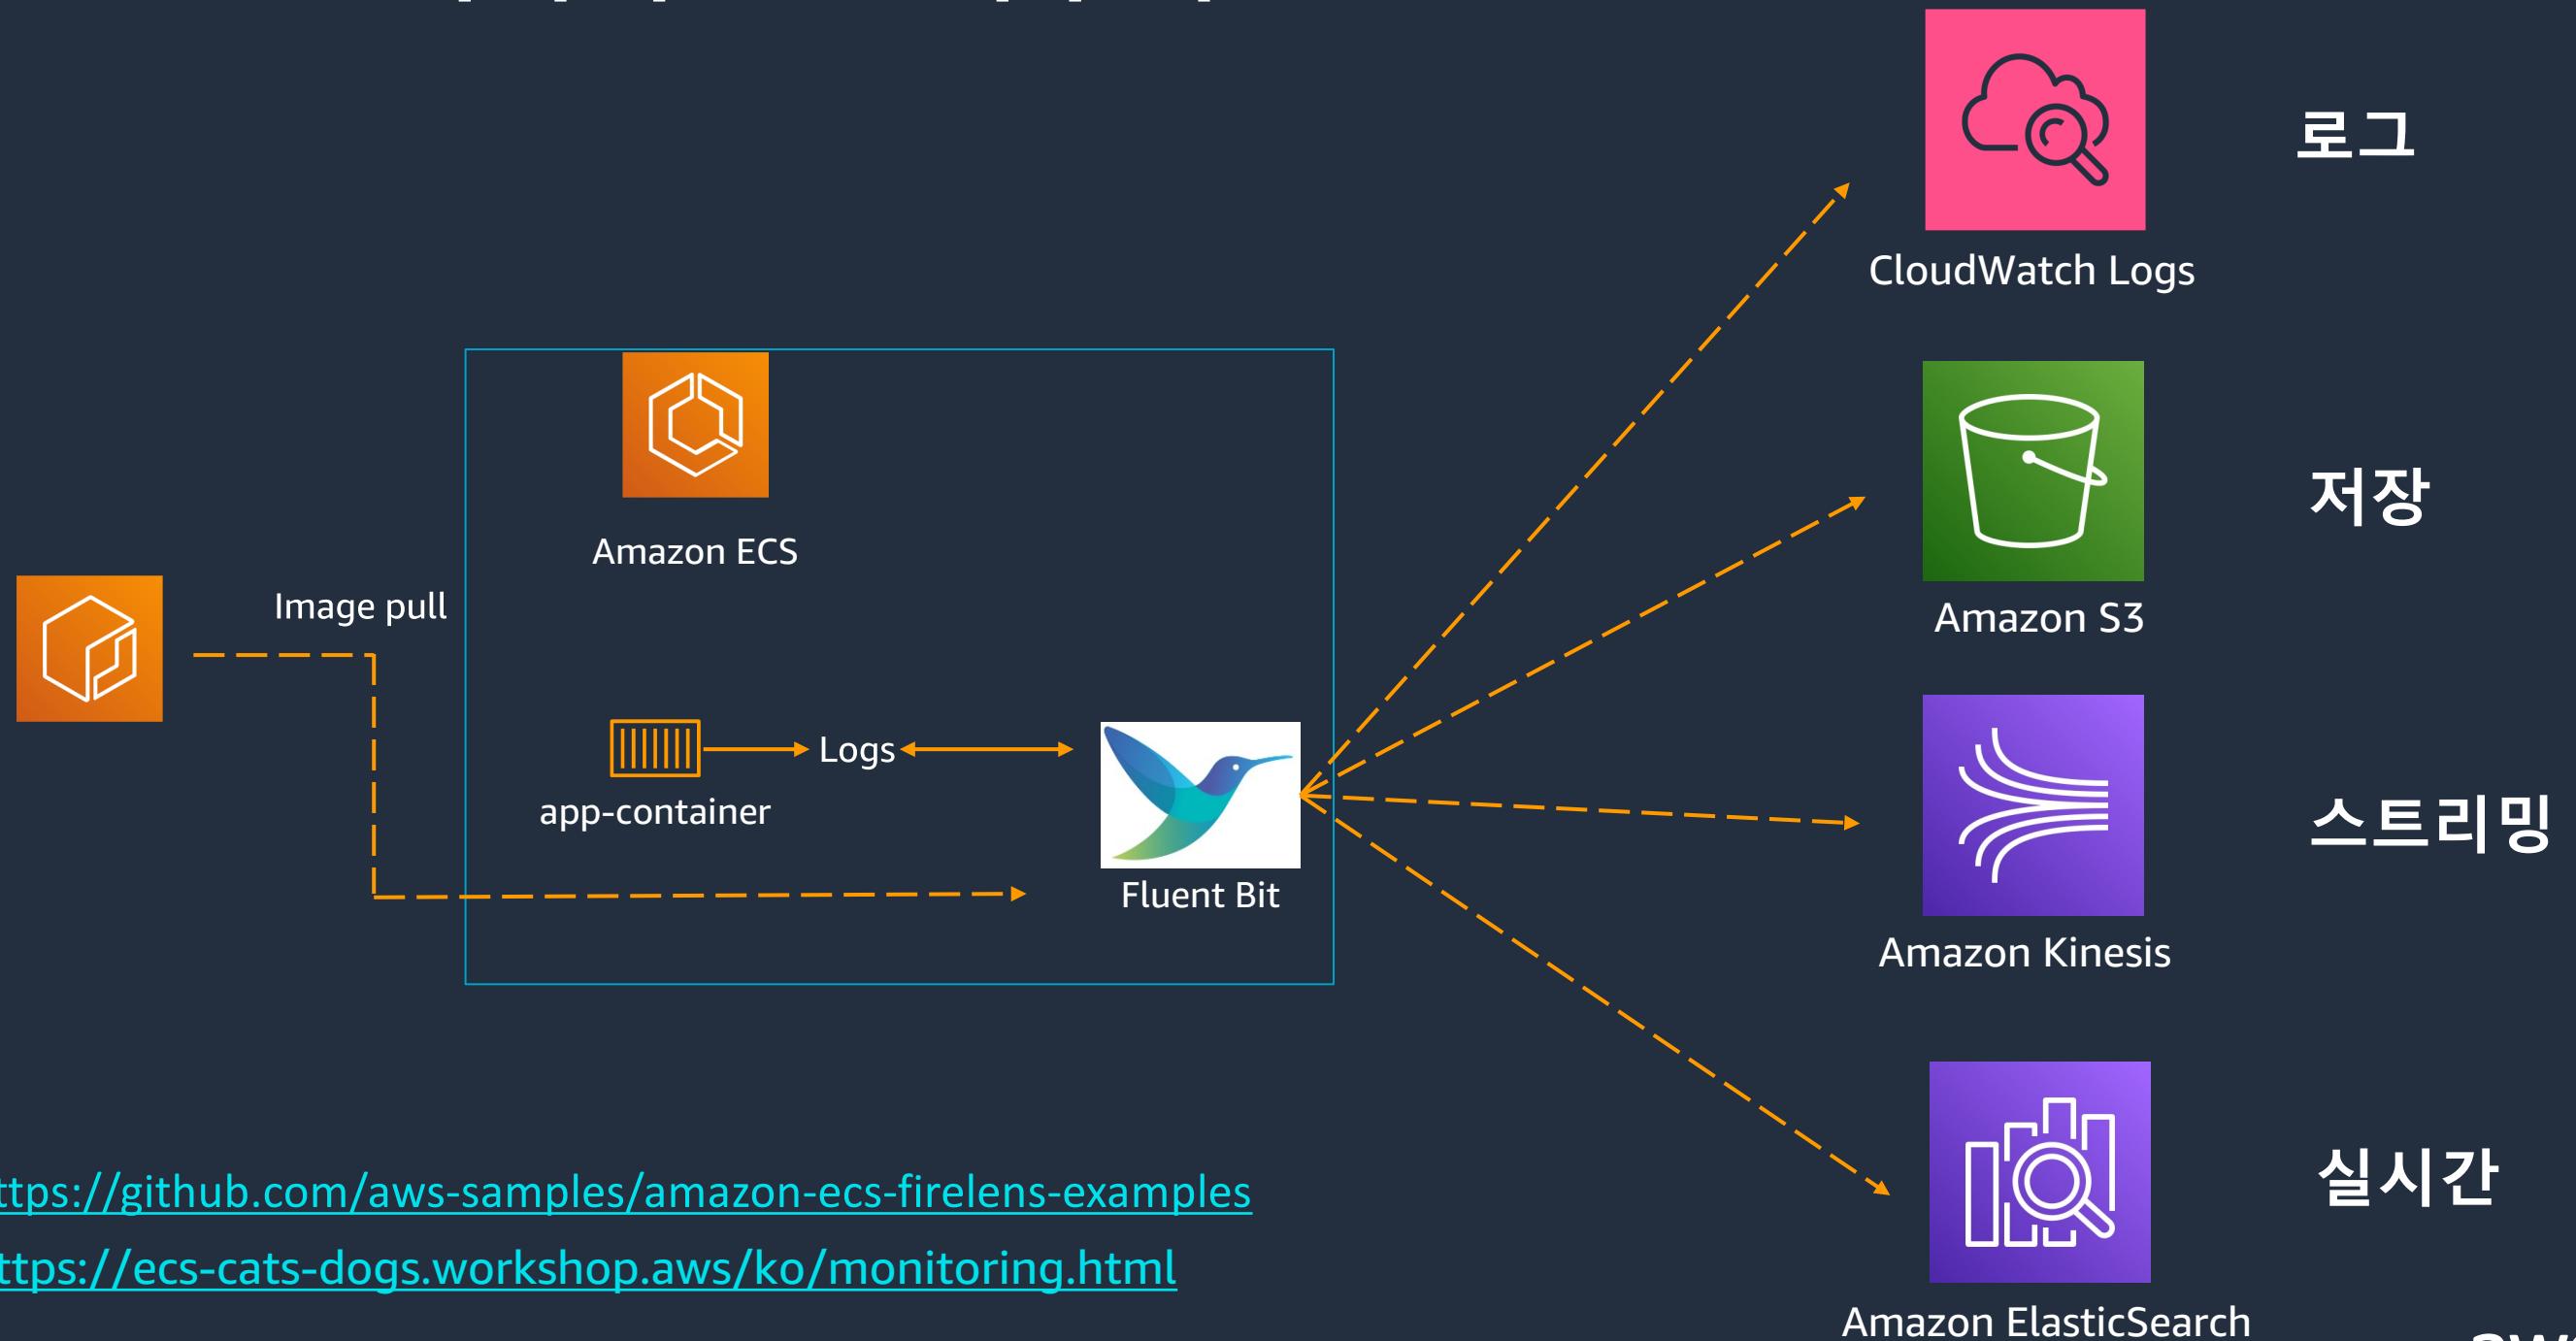
<https://ecs-cats-dogs.workshop.aws/ko/monitoring.html>

Amazon ECS CloudWatch Container Insights

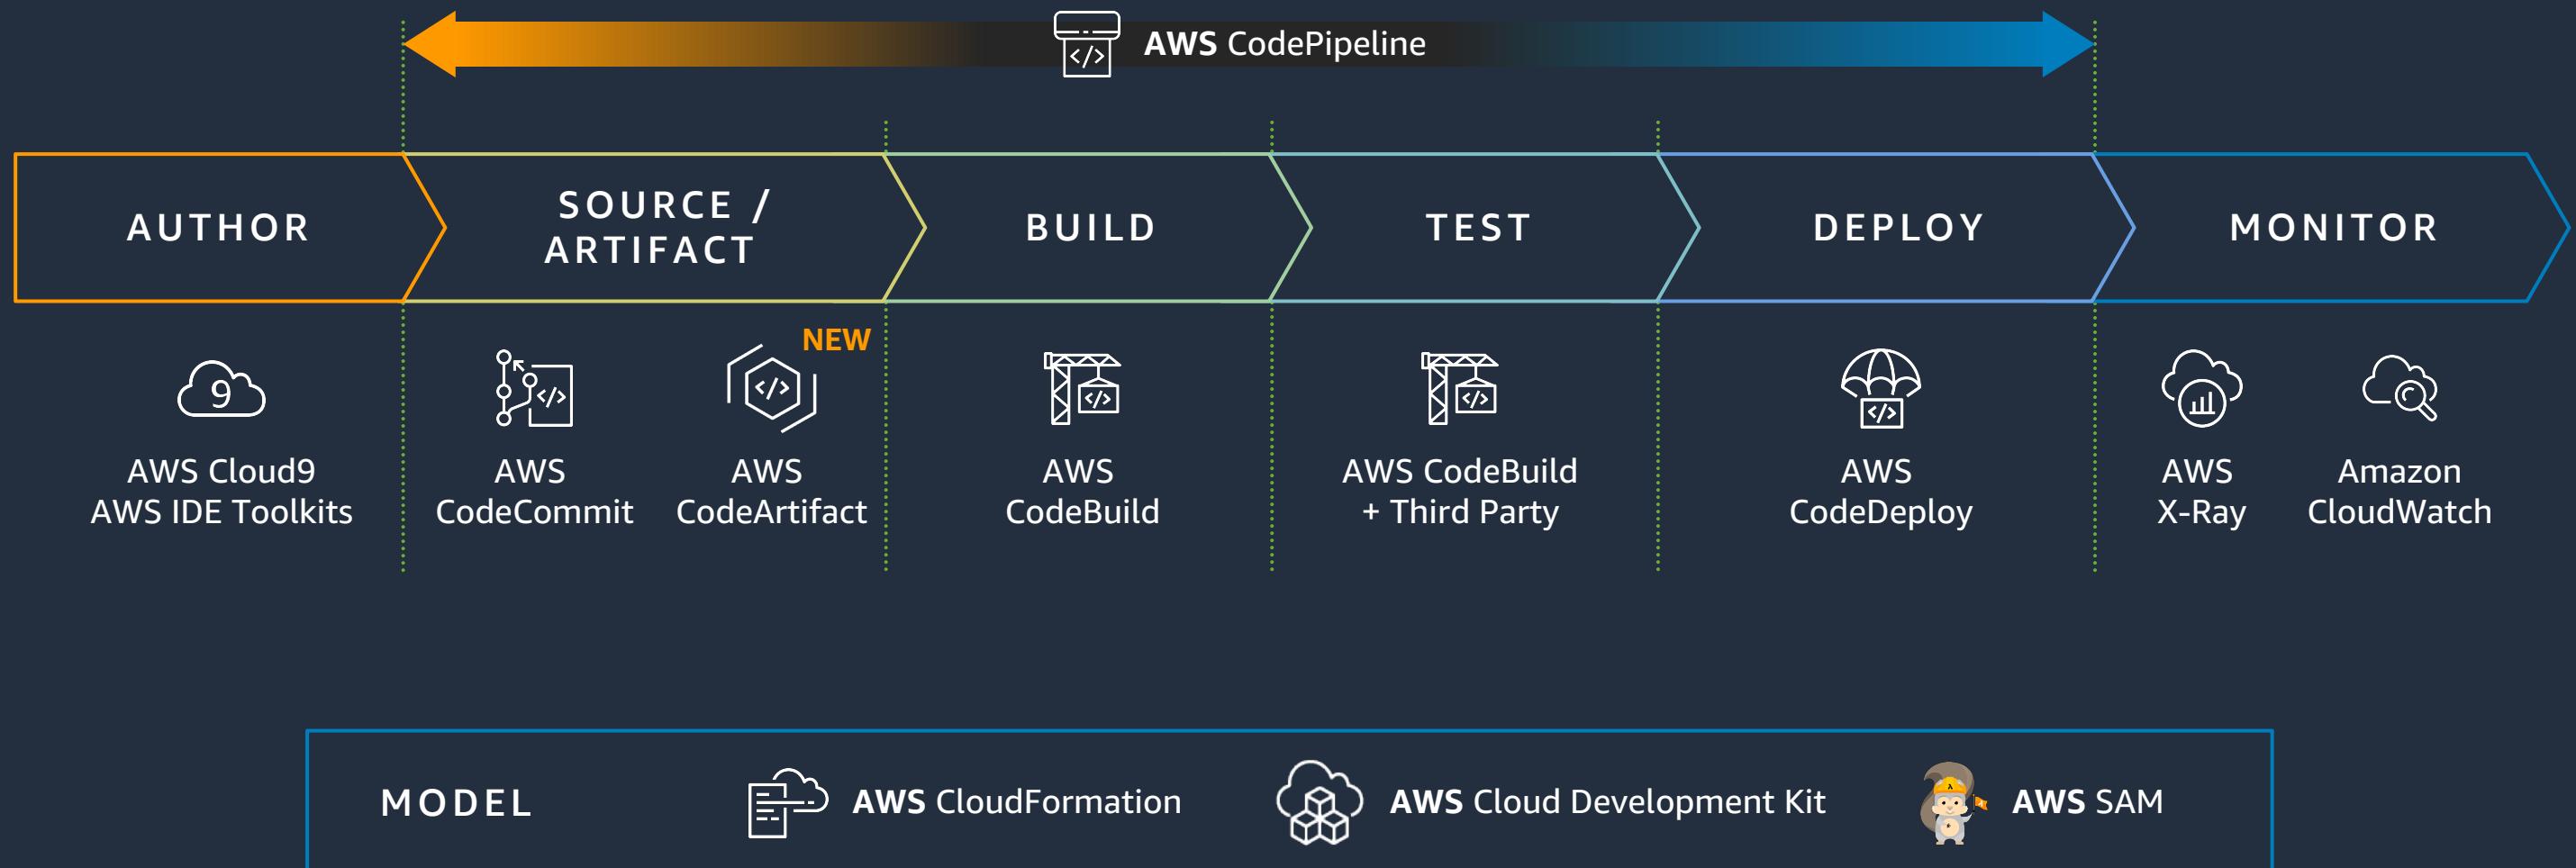


<https://ecs-cats-dogs.workshop.aws/ko/monitoring.html>

FireLens: 컨테이너 로그 라우터



AWS Developer Tool



<https://ecs-cats-dogs.workshop.aws/ko/cicd.html>

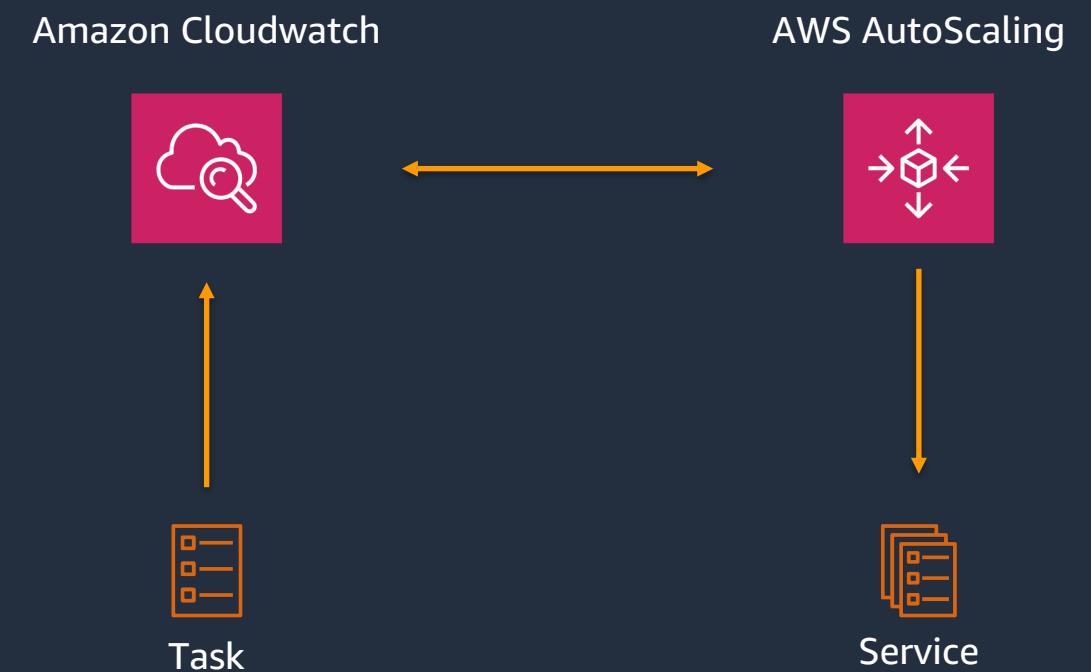
ECS Auto Scaling and Capacity Provider

<https://ecs-cats-dogs.workshop.aws/ko/autoscale.html>



ECS Service Auto Scaling

서비스 내의 작업 수를 자동으로 늘리거나
줄이는 역할을 함
ECS는 CPU 및 메모리 통계를 CloudWatch에
제시하고, Target Tracking, Step Scaling, 그리고
Scheduled Scaling을 지원함



<https://ecs-cats-dogs.workshop.aws/ko/autoscale/service.html>

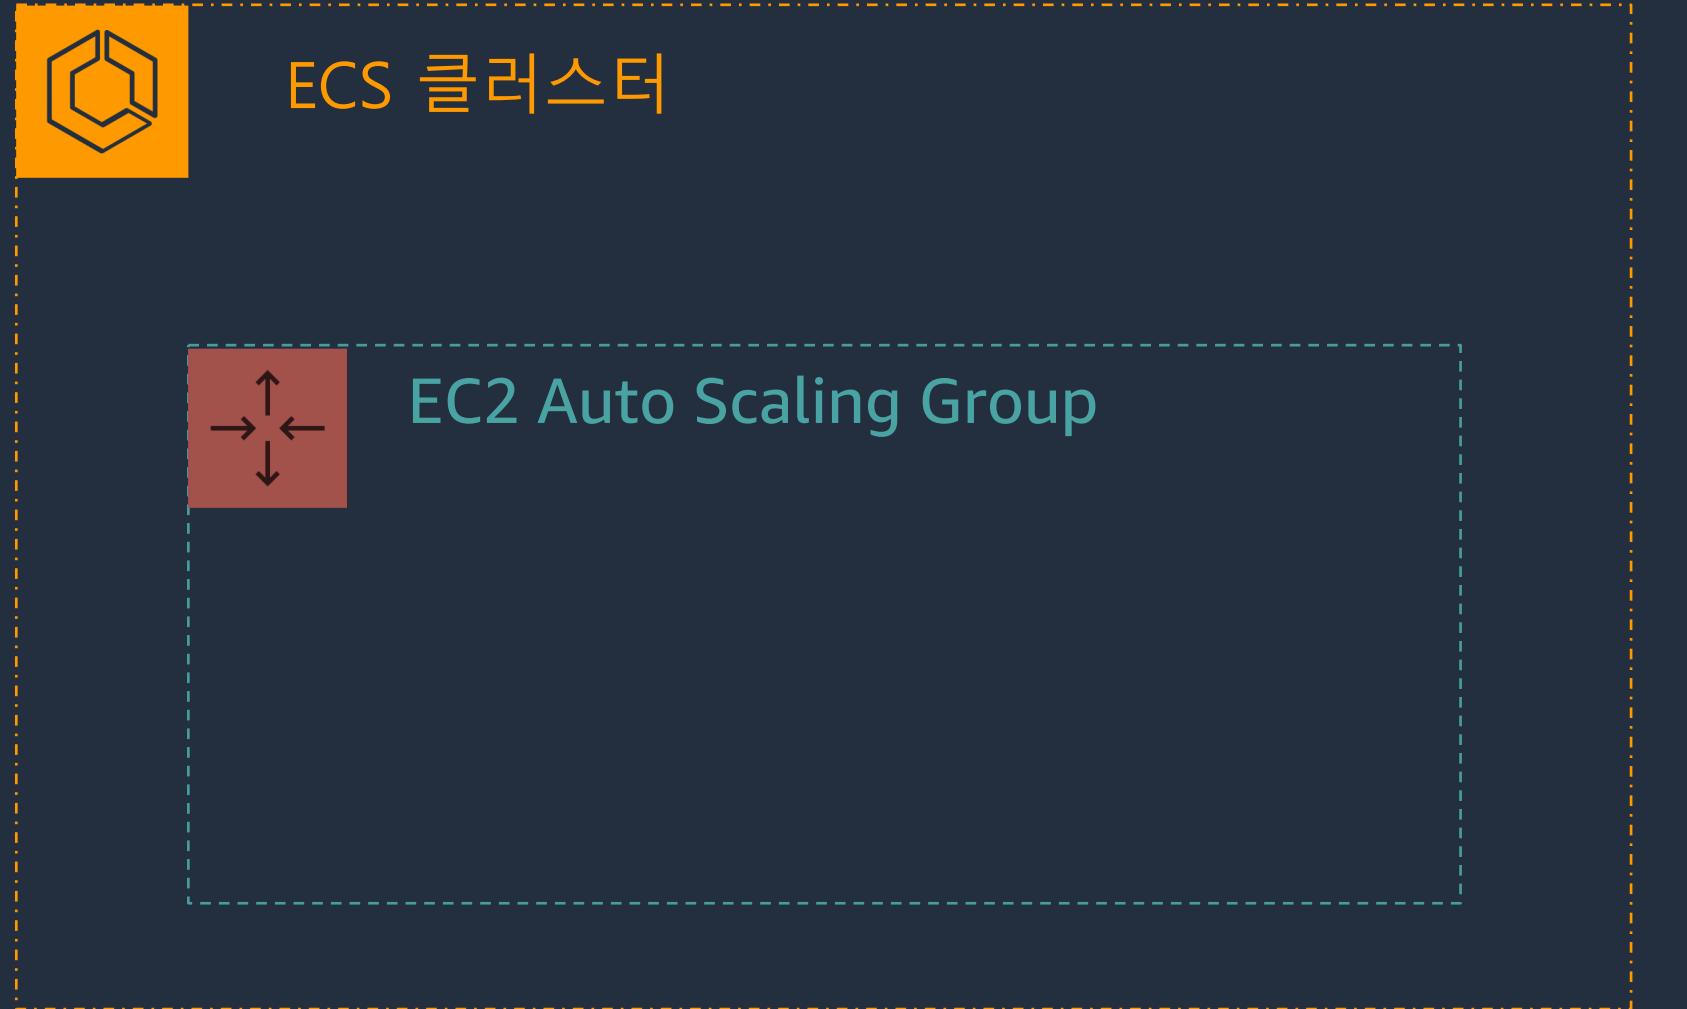
ECS 클러스터와 Amazon EC2 Auto Scaling 그룹



ECS 클러스터

클러스터 생성

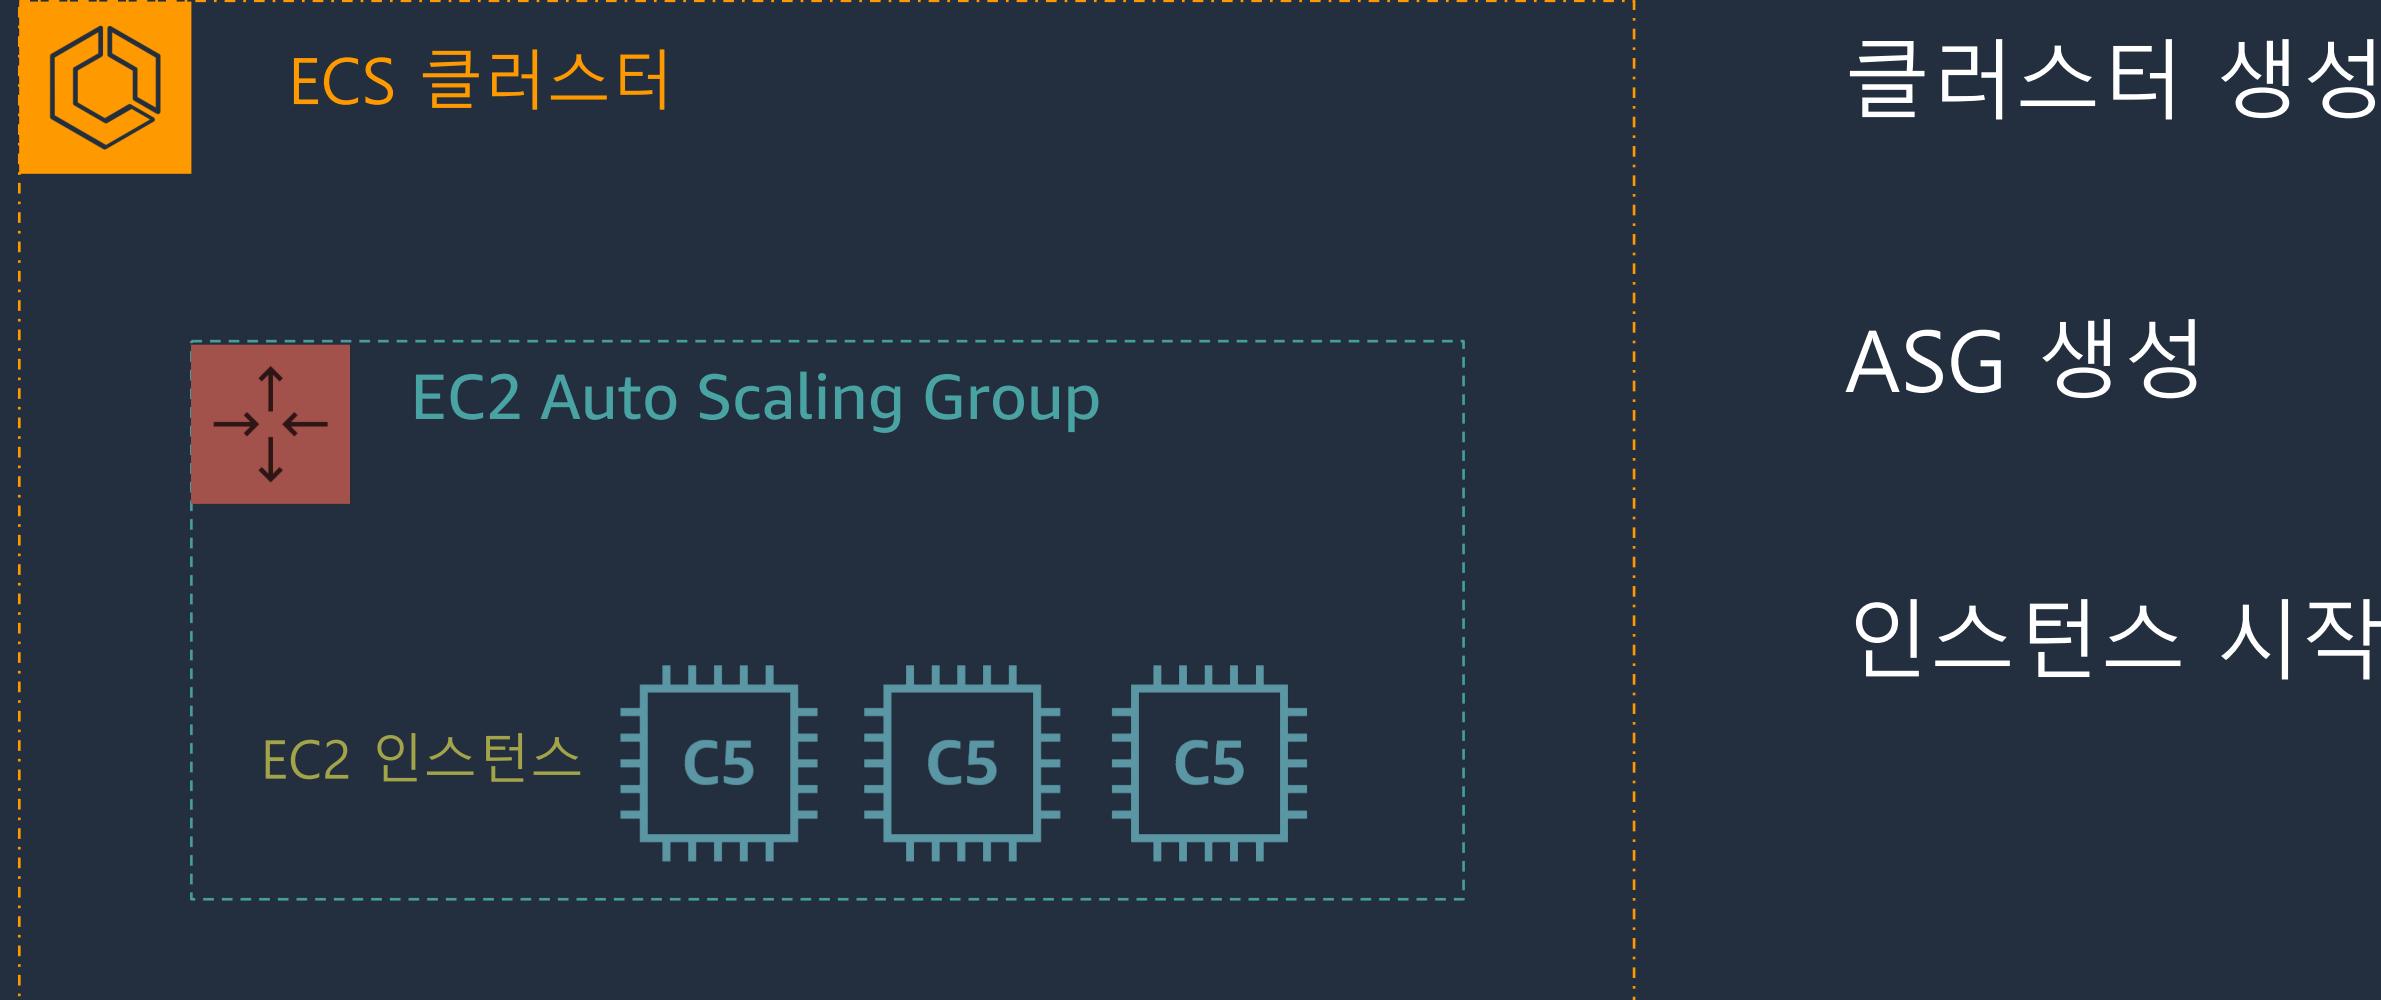
ECS 클러스터와 Amazon EC2 Auto Scaling 그룹



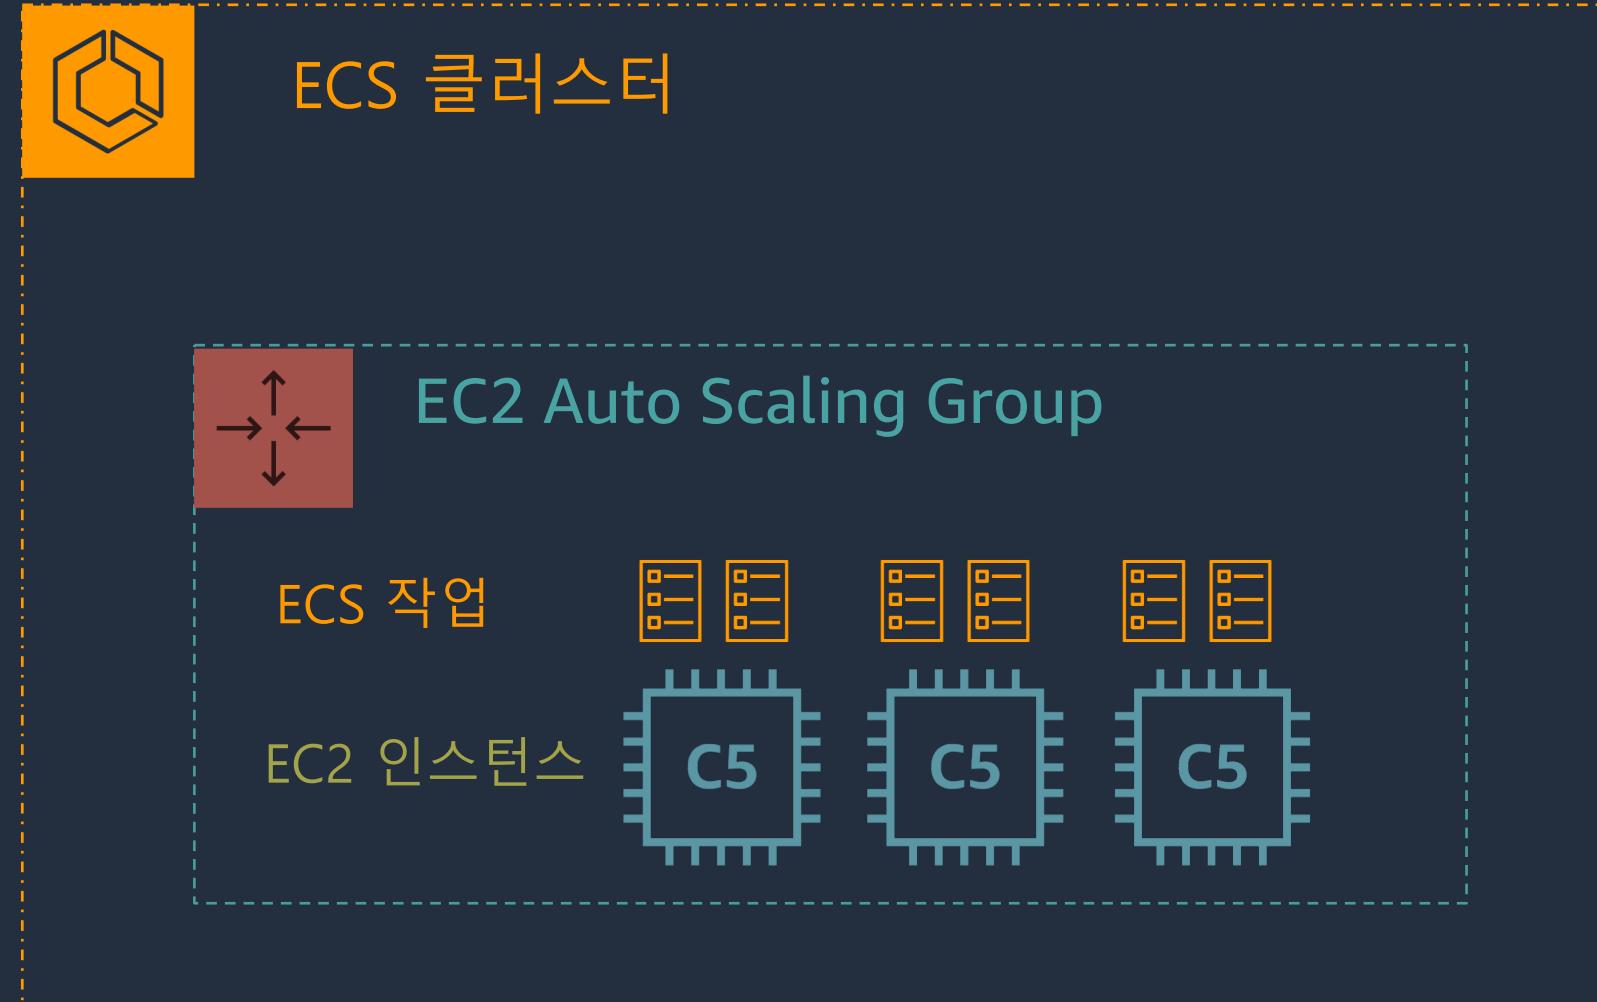
클러스터 생성

ASG 생성

ECS 클러스터와 Amazon EC2 Auto Scaling 그룹



ECS 클러스터와 Amazon EC2 Auto Scaling 그룹



클러스터 생성

ASG 생성

인스턴스 시작

작업 수행

애플리케이션 우선



애플리케이션 우선 주의

애플리케이션은 각자의 요구사항이 있음

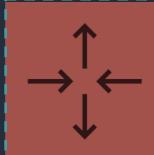
인프라스트럭처는 애플리케이션의 요구조건에 응답해야 함

Amazon ECS 용량 공급자: 연결고리



ECS 클러스터

ECS 용량 공급자



EC2 Auto Scaling Group

클러스터 생성 (1회)

ASG 생성 (1회)

용량 공급자 생성(1회)

Amazon ECS 용량 공급자: 연결고리



클러스터 생성 (1회)

ASG 생성 (1회)

용량 공급자 생성(1회)

작업 실행

Amazon ECS 용량 공급자: 연결고리



클러스터 생성 (1회)

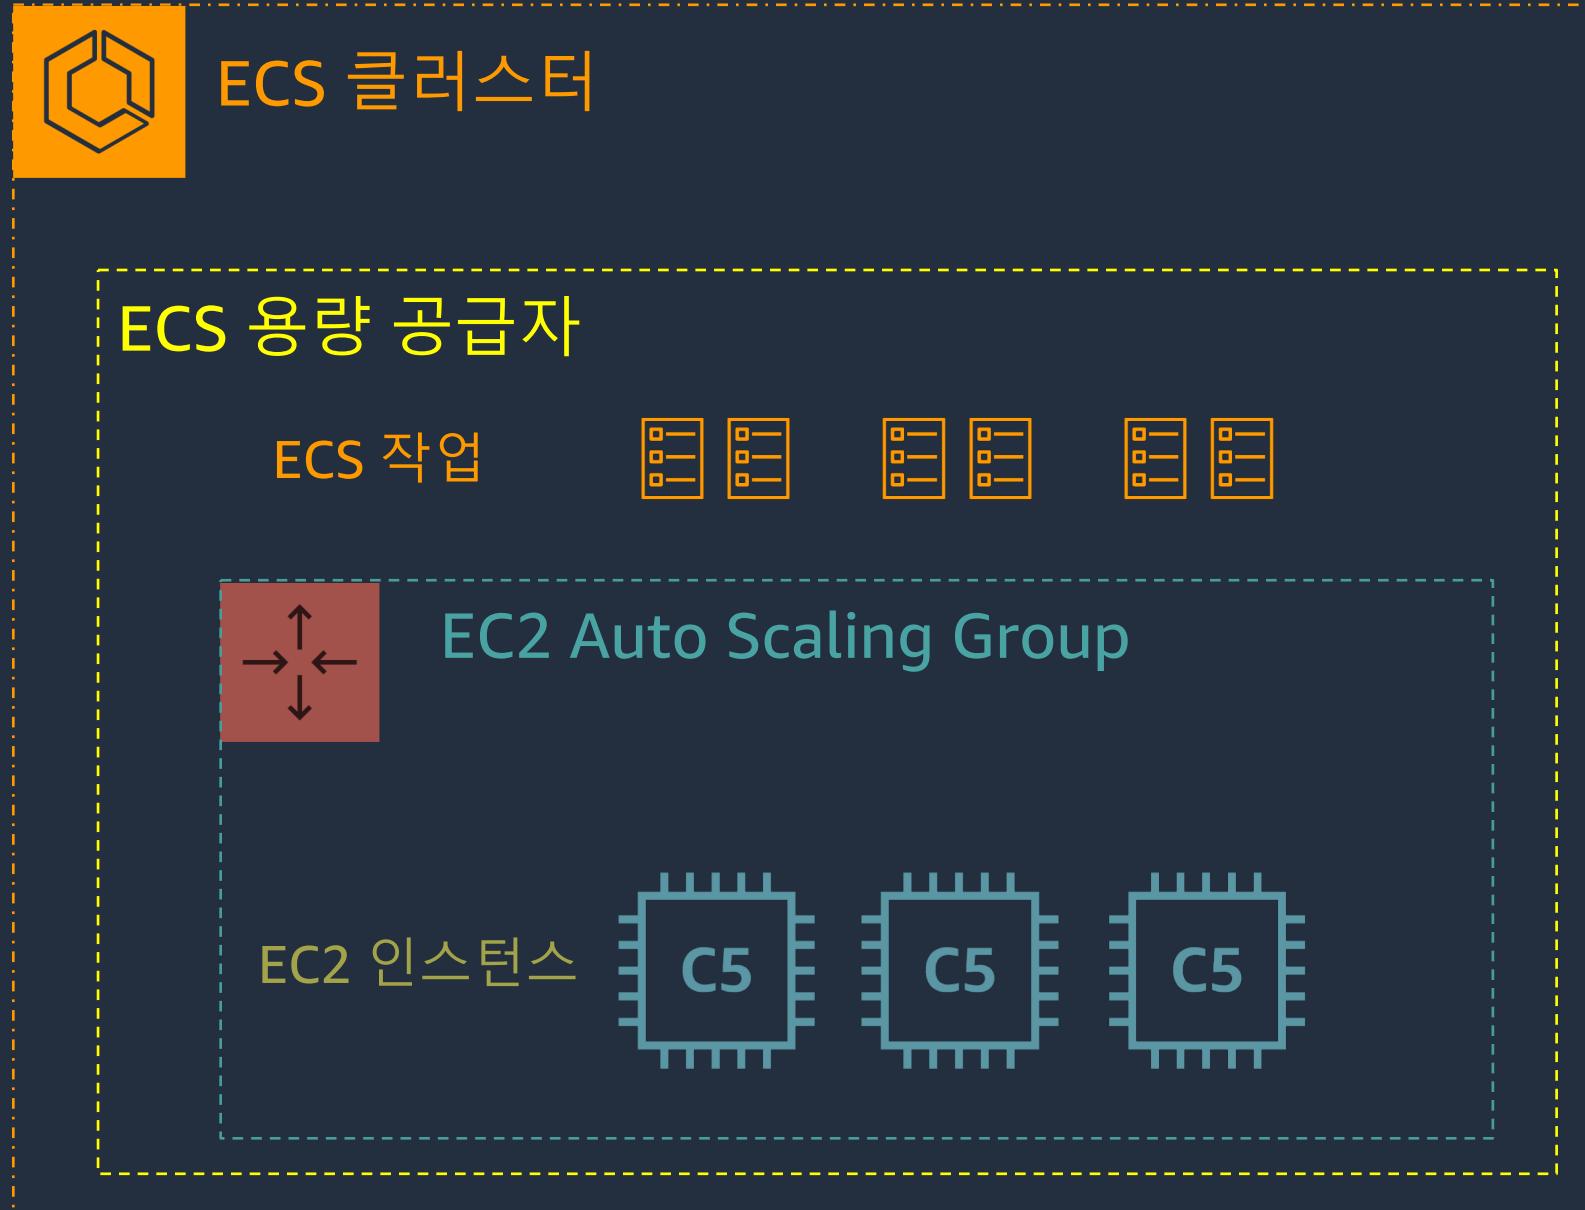
ASG 생성 (1회)

용량 공급자 생성(1회)

작업 실행

인스턴스 시작

Amazon ECS 용량 공급자: 연결고리



클러스터 생성 (1회)

ASG 생성 (1회)

용량 공급자 생성(1회)

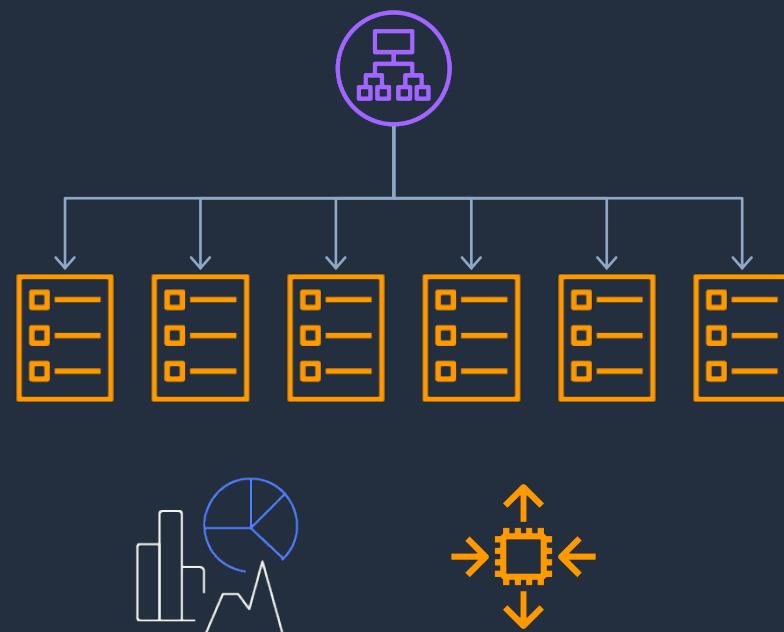
작업 실행

인스턴스 시작

작업 배치

여러 용량 공급자로 분할: On-demand와 Spot

50% 여유 리소스 준비:
2/3 on-demand, 1/3 spot

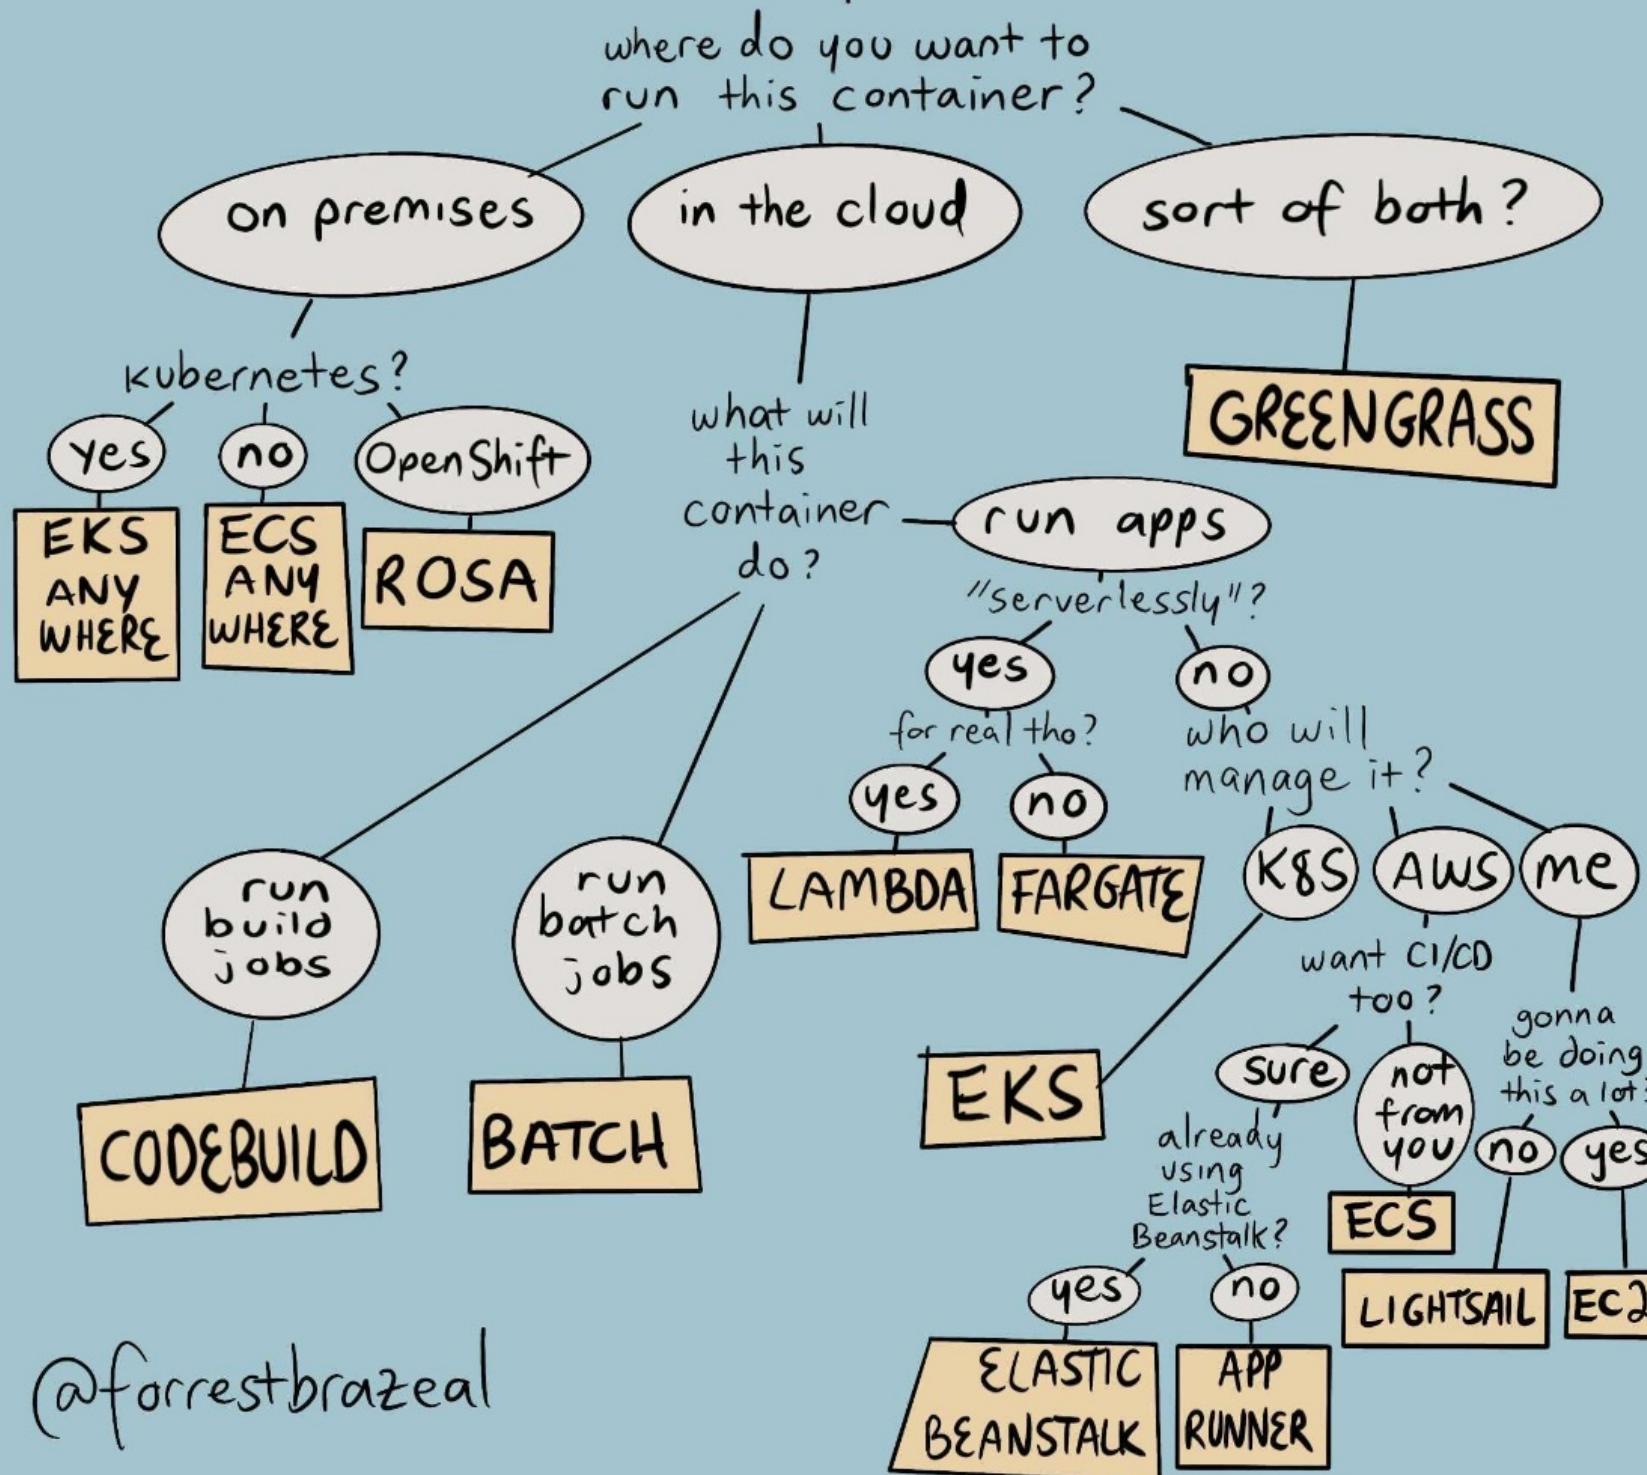


+50% 여유
+5-10% 추가 비용



<https://aws.amazon.com/ko/blogs/containers/deep-dive-on-amazon-ecs-cluster-auto-scaling/>

WHICH AWS CONTAINER SERVICE SHOULD I USE?



Start Containerization

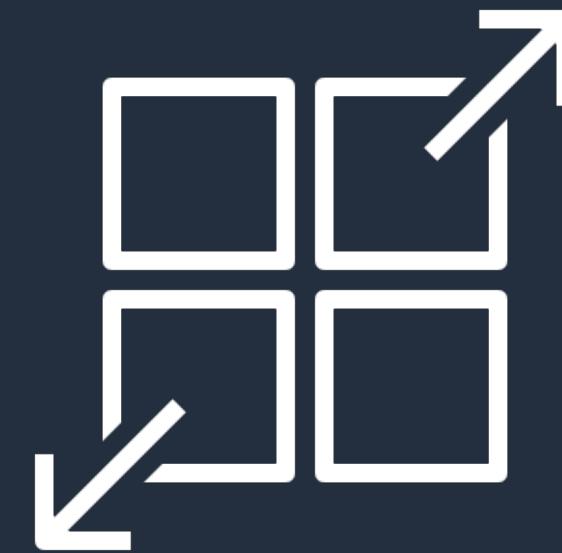


ECS Cats and Dogs 워크샵

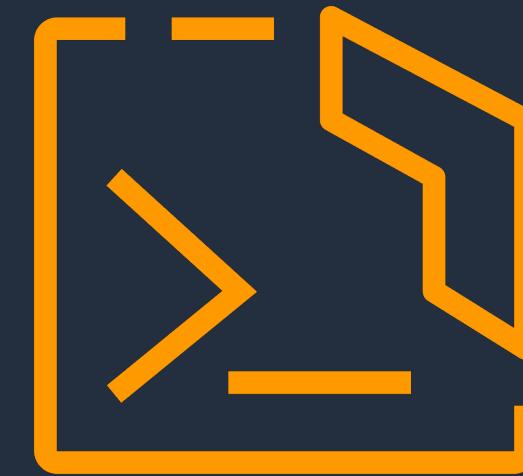
<https://ecs-cats-dogs.workshop.aws/ko/>



컨테이너화를 쉽게 도와줄 도구들



AWS App Runner



AWS Copilot CLI

* 6월 기준 서울 리전 미지원

AWS App Runner

The screenshot shows the AWS App Runner configuration interface. On the left, the 'Source' tab is selected, showing options for 'Repository type'. 'Container registry' is unselected, while 'Source code repository' is selected, highlighted with a blue border. Below this, there's a 'Connect to GitHub' section with a dropdown for 'rothgar' and an 'Add new' button. Under 'Repository', 'hello-app-runner' is selected from a dropdown, with a refresh icon next to it. Under 'Branch', 'main' is selected from a dropdown, also with a refresh icon. On the right, the 'Build settings' tab is selected. It includes sections for 'Configuration file' (with 'Configure all settings here' selected), 'Runtime' (set to 'Python 3'), 'Build command' (containing the command 'yum install -y pycairo && pip install -r requirements.txt'), 'Start command' (containing 'python app.py'), and 'Port' (set to '8000').

Source

Repository type

Container registry
Deploy your service from a container image stored in a container registry.

Source code repository
Deploy your service from code hosted in a source code repository.

Connect to GitHub [Info](#)

App Runner deploys your source code by installing an app called "AWS Connector for GitHub" in your account. You can install this app in your main GitHub account or in a GitHub organization.

rothgar

Add new

Repository

hello-app-runner

Branch

main

Build settings

Configuration file

Configure all settings here
Specify all settings for your service here in the App Runner console.

Use a configuration file
Let App Runner read your configuration from the apprunner.yaml file in your source repository.

Runtime

Choose an App Runner runtime for your service.

Python 3

Build command

This command runs in the root directory of your repository when a new code version is deployed. Use it to install dependencies or compile your code.

```
yum install -y pycairo && pip install -r requirements.txt
```

Start command

This command runs in the root directory of your service to start the service processes. Use it to start a webserver for your service. The command can access environment variables that App Runner and you defined.

```
python app.py
```

Port

Your service uses this IP port.

8000

<https://aws.amazon.com/ko/blogs/korea/app-runner-from-code-to-scalable-secure-web-apps/>

Copilot Resources

- Environment
- Application
- Service
- Task
- Pipeline
- Storage

<https://aws.amazon.com/ko/blogs/containers/introducing-aws-copilot/>

```
Commands
Getting Started 🌱
  init      Create a new ECS application.
  docs      Open the copilot docs.

[

Develop ✨
  app      Commands for applications.
           Applications are a collection of services and environments.

  env      Commands for environments.
           Environments are deployment stages shared between services.

  svc      Commands for services.
           Services are long-running Amazon ECS services.

  job      Commands for jobs.
           Jobs are tasks that are triggered by events.

  task      Commands for tasks.
           One-off Amazon ECS tasks that terminate once their work is done.

Release 🚀
  pipeline  Commands for pipelines.
           Continuous delivery pipelines to release services.

  deploy    Deploy a Copilot job or service.

Addons 🐾
  storage   Commands for working with storage and databases.

Settings 🛡️
  version   Print the version number.
  completion Output shell completion code.

Flags
  -h, --help    help for copilot
  -v, --version version for copilot

Examples
  Displays the help menu for the "init" command.
  '$ copilot init --help'
```



감사합니다 😊

jimini@amazon.com

